

Operating system

Part IV: CPU Scheduling algorithms

By KONG LingBo (孔令波)

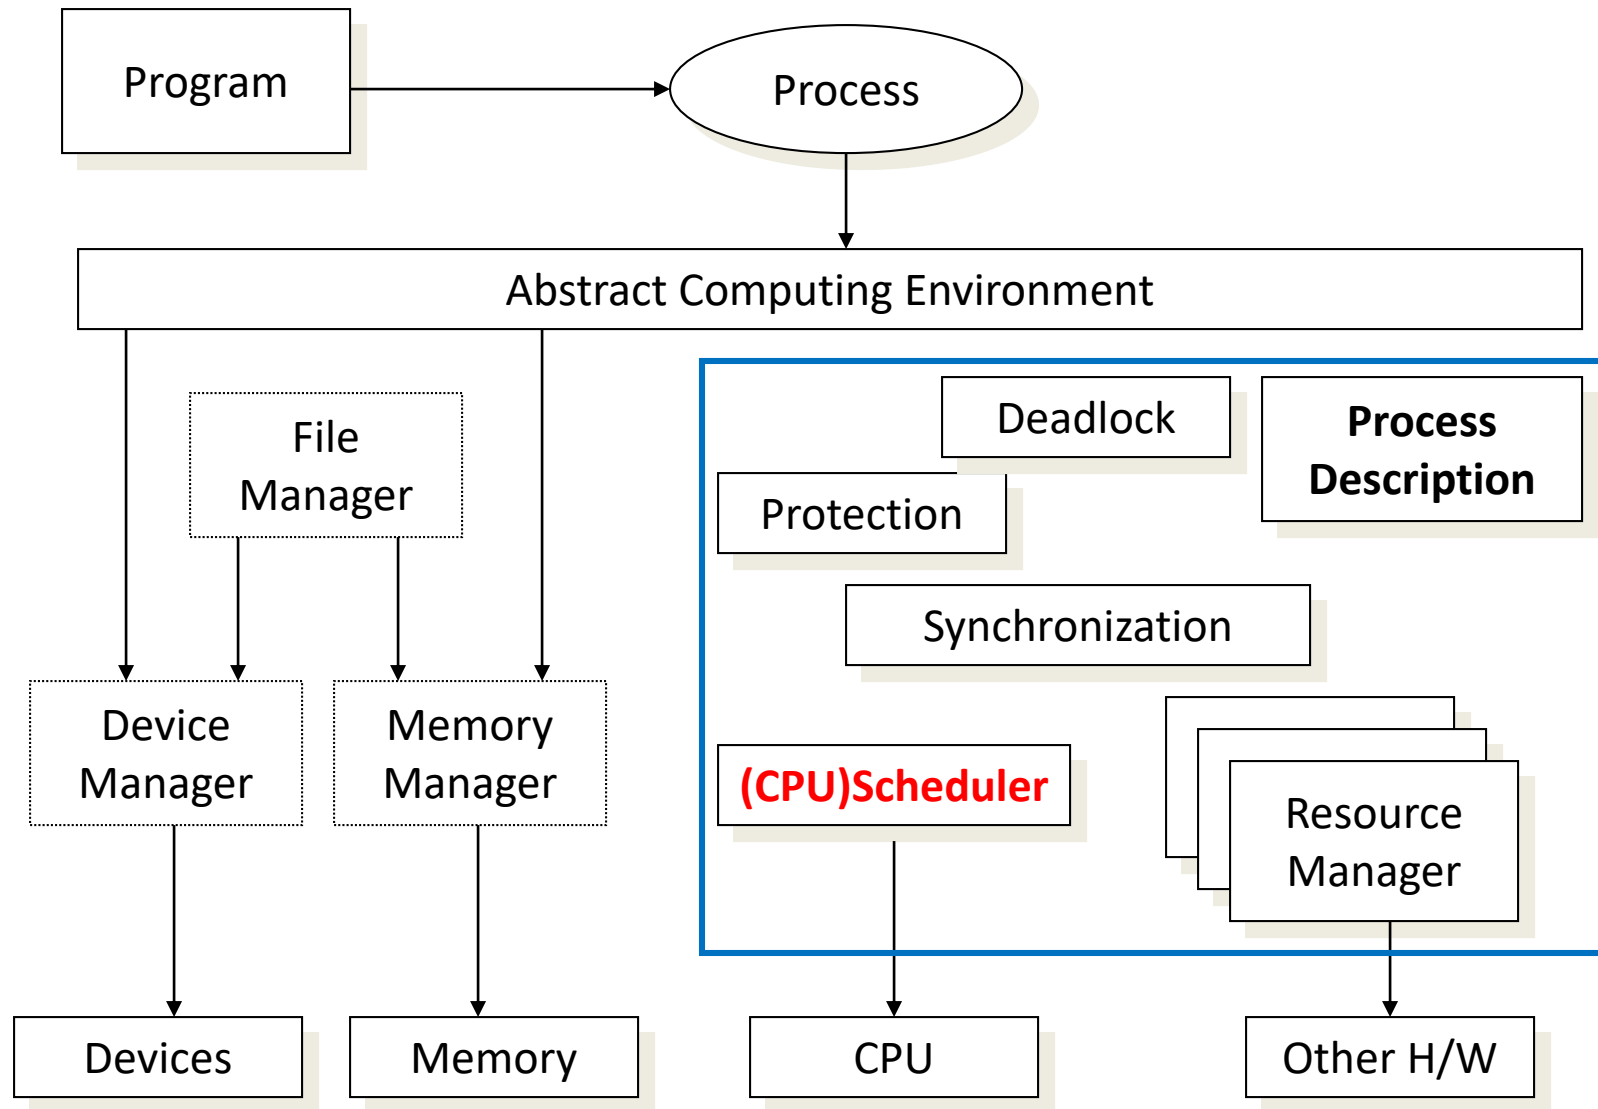
Outline of topics covered by this course

- **INTRODUCTION** – why should we learn OS?
- **OVERVIEW** – what problems are considered by modern OS in more details?
- **EXECUTION – CPU management**
 - Process and Thread
 - CPU scheduling
- **EXECUTION – competition (synchronization problem)**
 - Synchronization
 - Deadlock


Goals

- To know the mechanism of CPU scheduling
 - Basic Concepts, Criteria
- CPU scheduling Algorithms
 - FCFS, SJF (Non-Preemptive), SRJF (Preemptive), Priority, Round Robin, Multilevel Queue, Multilevel Feedback Queue, Lottery etc.
- Algorithms Evaluation
 - How to simulate and evaluate those CPU scheduling algorithms

Overview of OS



Three kinds of schedulers



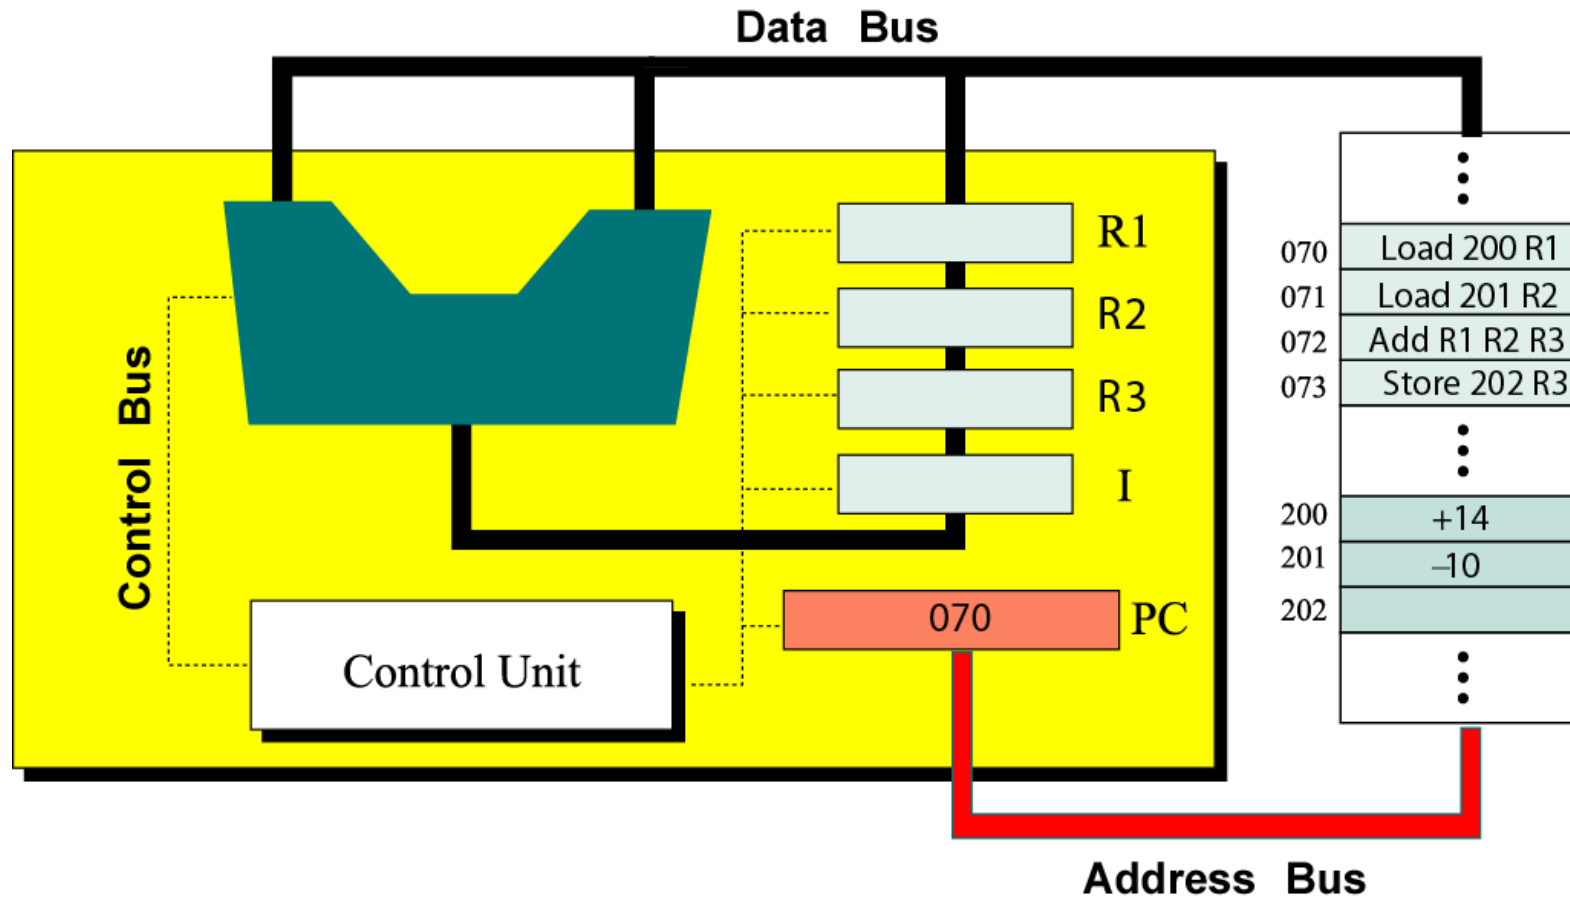
Short-term scheduler

determines which **processes** could get the usage of CPU

- Basic Concepts
 - CPU and some IO devices could work in parallel
 - Process could be categorized into IO-burst and CPU burst types
- Scheduling Criteria & Metrics
- Different Scheduling Algorithms
- Algorithm Evaluation

We have learned

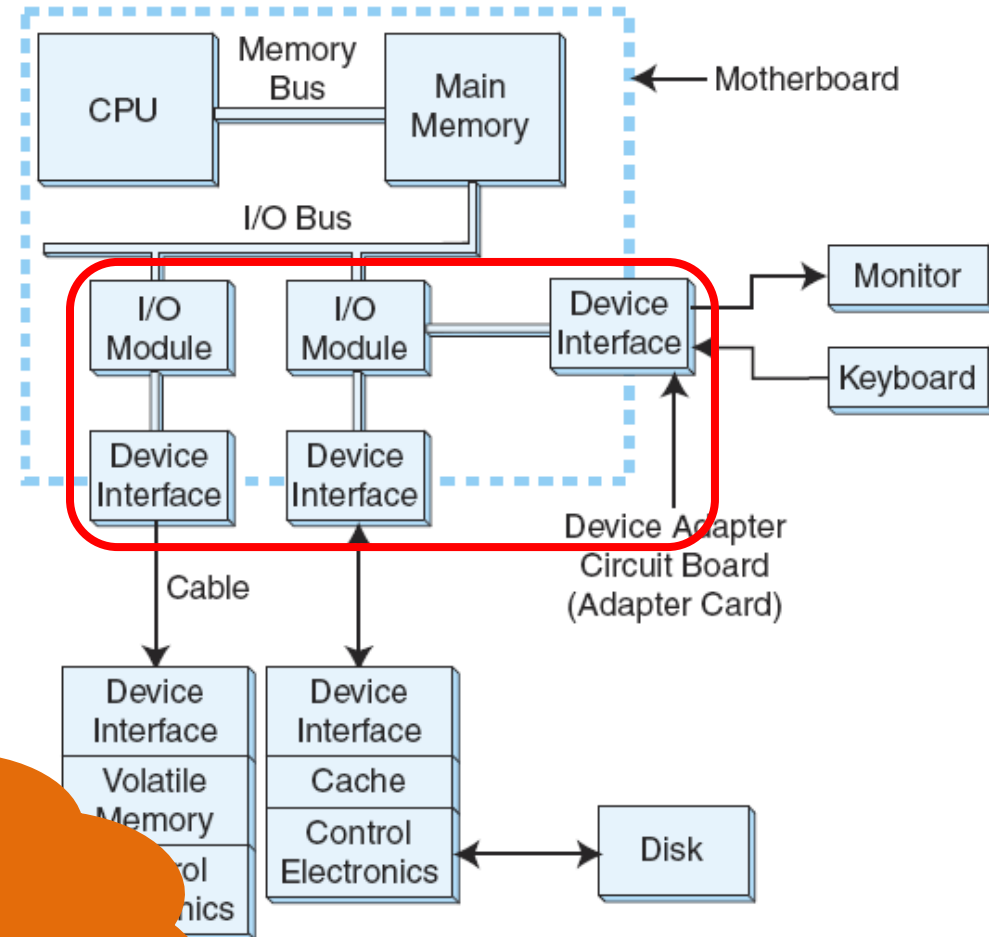
- CPU is the most important resource
 - Its job is to execute the instructions following Machine Cycle



Now

- CPU and IO can work in parallel
 - CPU now focuses on the **computation**
 - IO chipsets take over the IO task

We'll learn the framework to manage IO devices in later IO chapter



1.1 A Model I/O Configuration

[2003.Essentials Of Computer Organization And Architecture+++++.pdf]

Now

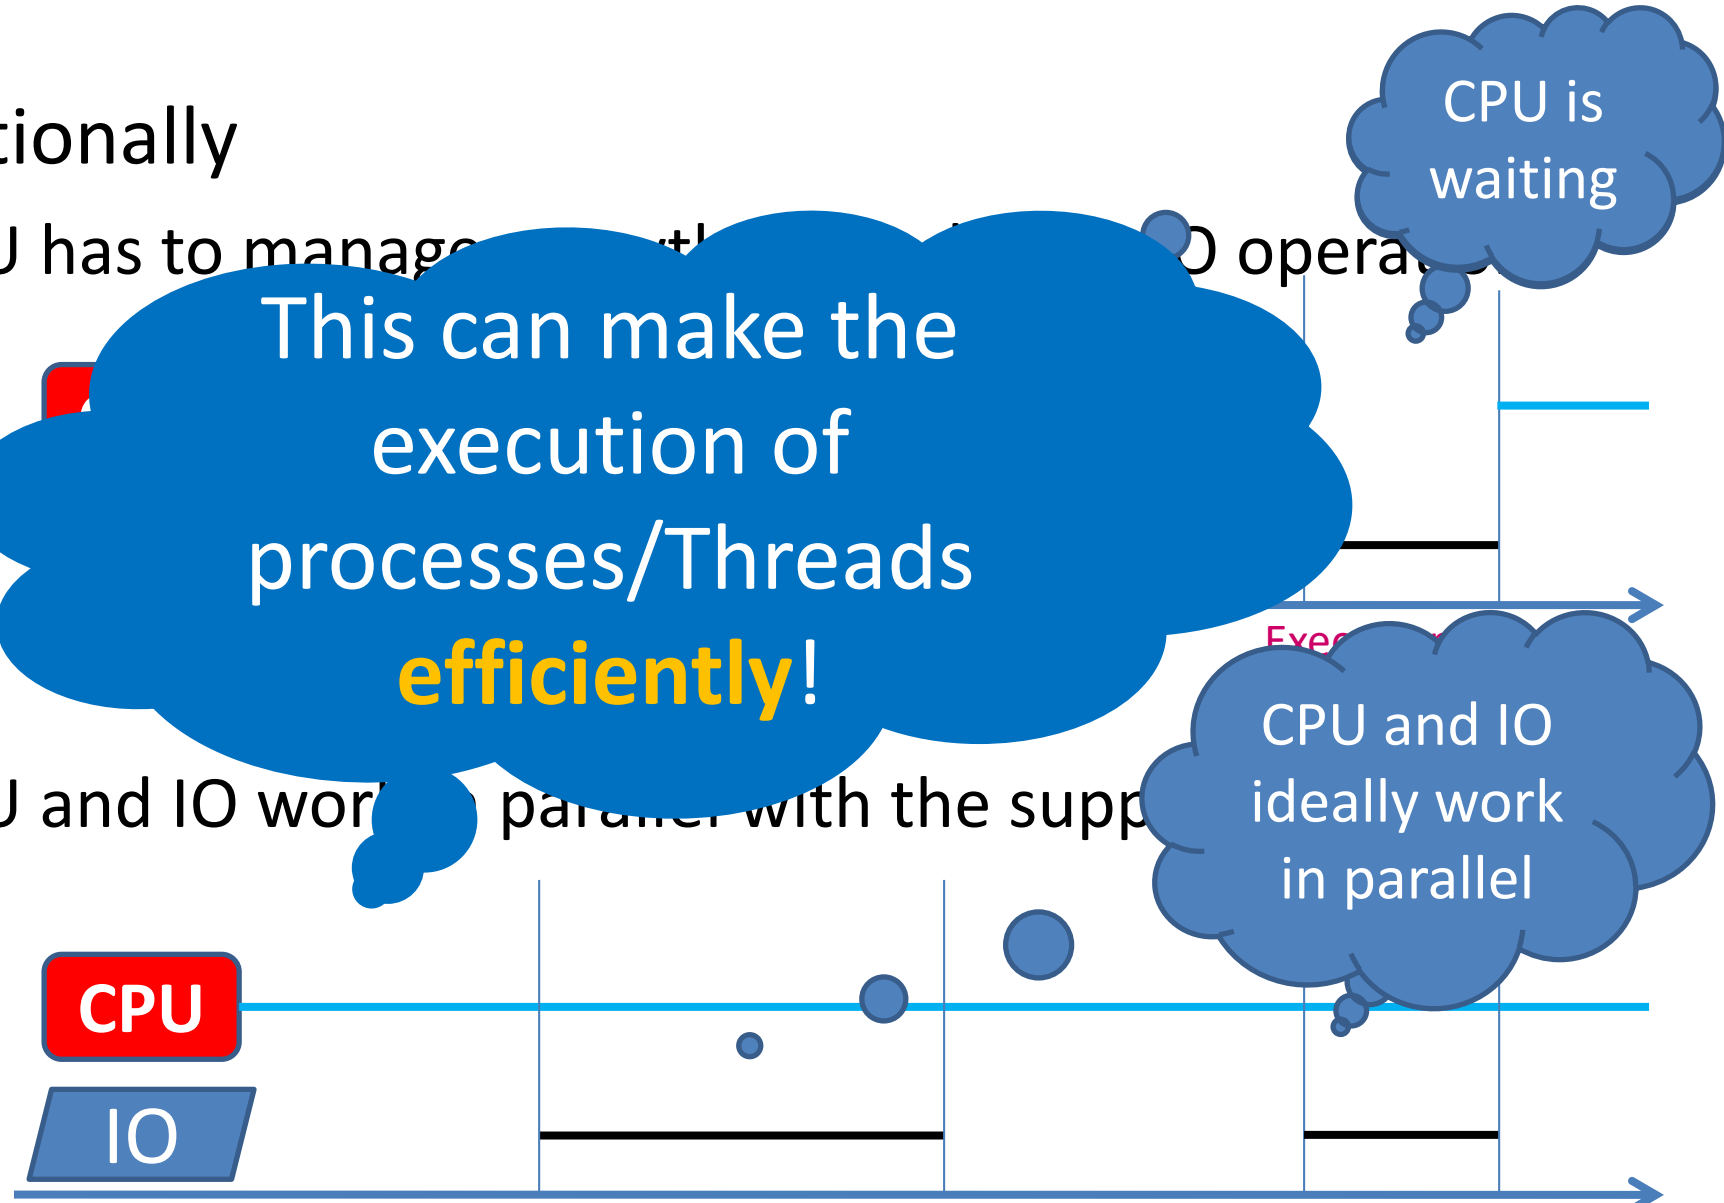
- Traditionally

- CPU has to manage with the I/O operation

This can make the
execution of
processes/Threads
efficiently!

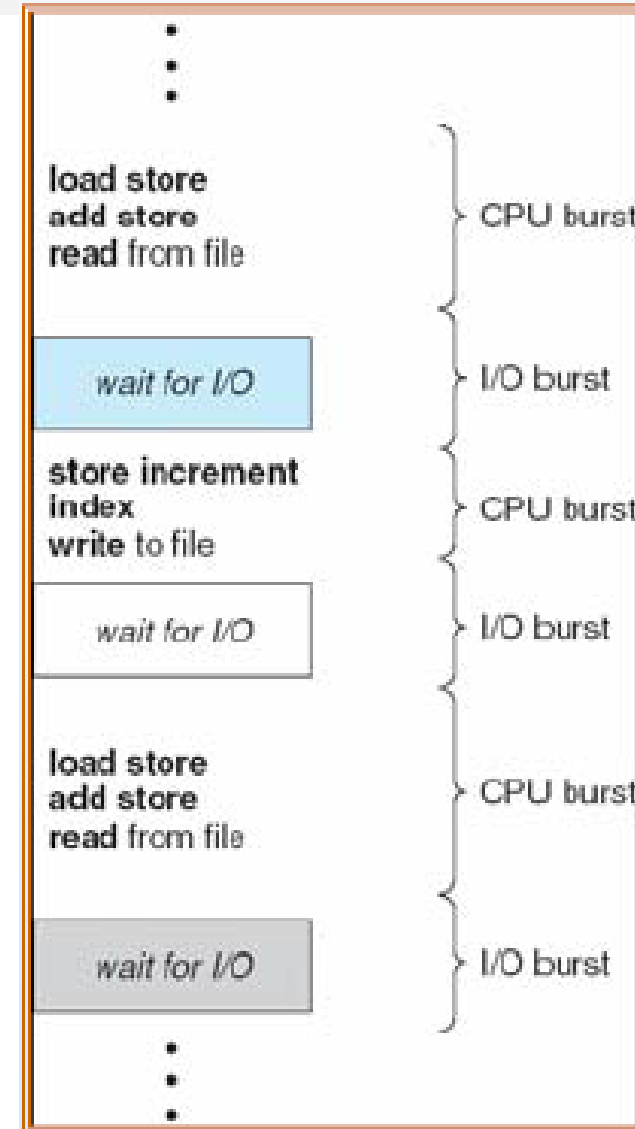
- Now

- CPU and IO work in parallel with the support of



Process's property: CPU-I/O Burst [突发, 爆发] Cycle

- Process execution consists of a cycle of **CPU execution** and **I/O request**
- Process execution begins with a **CPU burst** and followed by an **I/O burst**



CPU-I/O Burst


- Burst/Service time = total processor time needed in one CPU-I/O burst cycle.
- Jobs/Process with long CPU burst time are **CPU-bound jobs/processes** and are also referred to as “long jobs/processes”.
- Jobs with short CPU burst time are **IO-bound jobs/processes** and are also referred to as “short jobs/processes”.
- CPU-bound processes have longer CPU bursts than I/O-bound processes.

- Basic Concepts
- Scheduling Criteria & Metrics
- Different Scheduling Algorithms
- Algorithm Evaluation

Parameters to evaluate the scheduling


- CPU utilization [**CPU使用率**] (**Efficiency**)
 - keep the CPU as busy as possible (from 0% to 100%)
- **Fairness**: each process gets a “fair share” of the CPU
- Throughput [**吞吐量**]
 - # of processes that complete their execution per time unit
- Turnaround time [**周转时间**]
 - amount of time to execute a particular Process
 - i.e. execution time + waiting time
- Waiting time [**等待时间**]
 - amount of time a process has been waiting in the ready queue
- Response time [**响应时间**]
 - amount of time it takes from when a request was submitted until the first response is produced , not output (for time-sharing environment)

CPU utilization



- CPU Utilization
= $(10 + 30 + 25)/100$
= 65%

Turnaround time & Waiting time



- Turnaround time (tt_i) • Average Turnaround time (att)

= 100

$$att = \frac{\sum tt_i}{n}$$

- Waiting time (wt_i)

= 30 + 5


= 35

- Average Waiting time (awt)

$$awt = \frac{\sum wt_i}{n}$$


Goals for a Scheduler

- Suppose we have processes A, B, and C, submitted at time 0
- We want to know the response time, waiting time, and turnaround time of process **A**




Goals for a Scheduler

- Suppose we have processes A, B, and C, submitted at time 0
- We want to know the response time, waiting time, and turnaround time of process **B**



Goals for a Scheduler

- Suppose we have processes A, B, and C, submitted at time 0
- We want to know the response time, waiting time, and turnaround time of process **C**



You can derive the Optimization Criteria

- To maximize or minimize some average measures:
 - Max **CPU utilization**
 - Fraction of the time the CPU isn't idle
 - Max **throughput**
 - Amount of “useful work” done per time unit
 - Min **turnaround time**
 - Time from process creation to process completion
 - Min **waiting time**
 - Amount of time a process spends in the WAITING state
 - Min **Response time**

PPTs from others\OS PPT in English\ch06.ppt

Preemptive vs. non-preemptive scheduling

— 抢占式 v.s 非抢占式 调度

- **Non-preemptive** (not forcible removable):
 - Once a process is in the running state, it will continue until it terminates or blocks itself for I/O.
- **Preemptive** (forcible removable) :
 - Currently running process may be interrupted and moved to the Ready state by the OS.
 - processes can be suspended by scheduler
 - incurs a cost associated with access to share data.

Scheduling Algorithms

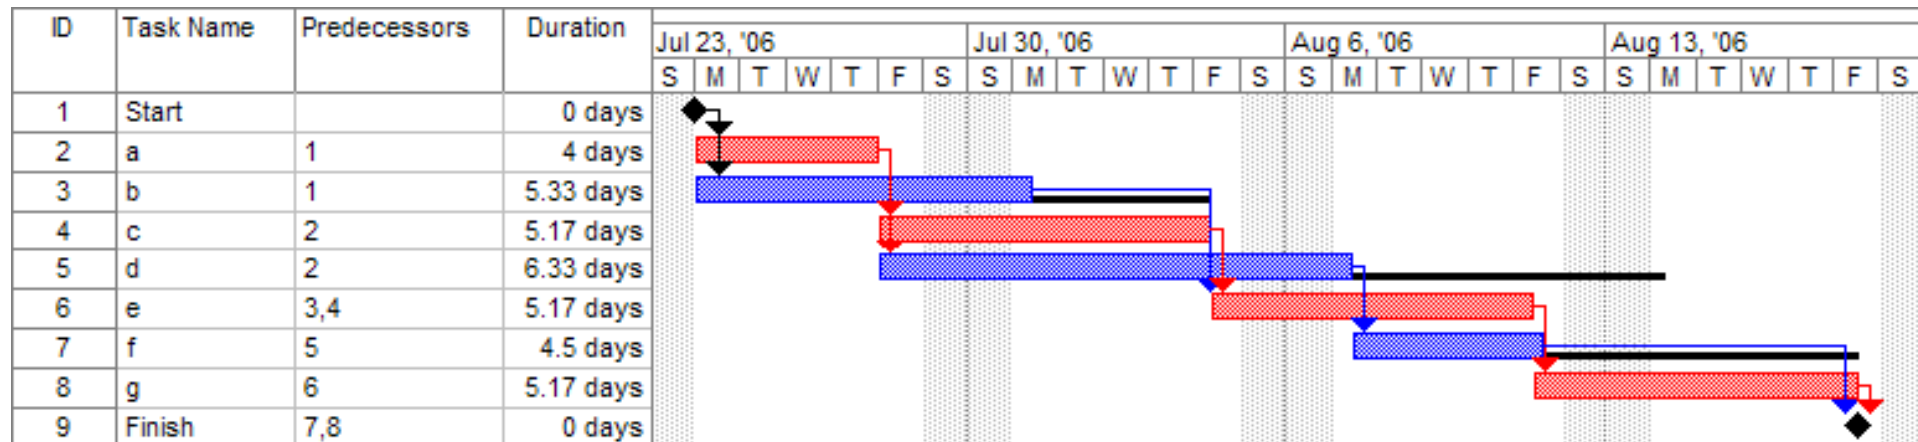
- **First Come First Serve Scheduling** [先来先服务]
(Non-preemptive)
- **Shortest Job First Scheduling** [最短任务先服务]
 - SRTF (Shortest Remaining Time First Scheduling)/SRJF
- **Priority Scheduling** [优先权]
- **Round-Robin Scheduling** [时间片轮转]
- **Multilevel Queue Scheduling** [多层次队列]
 - Multilevel Feedback-Queue Scheduling [多层次反馈队列]
- **Lottery Scheduling** [抽彩]

- Basic Concepts
- Scheduling Criteria & Metrics
- Different Scheduling Algorithms
- Algorithm Evaluation

Gantt Chart [甘特表]

http://en.wikipedia.org/wiki/Gantt_chart

- A Gantt chart is a type of bar chart, developed by Henry Gantt, that illustrates a project schedule.
 - Gantt charts illustrate the start and finish dates of the terminal elements and summary elements of a project.
 - Terminal elements and summary elements comprise the work breakdown structure of the project



FCFS (First Come First Serve) Scheduling


- FIFO scheduling
 - **Simplest** scheme
 - Processes dispatched according to arrival time
 - Non-preemptible/preemptive
 - Rarely used as primary scheduling algorithm
- Or First-In-First-Out (FIFO)



A Simple FCFS Example

<u>Process</u>	<u>Burst/CPU Time</u>	<u>Arrival time</u>
P_1	24	1
P_2	3	4
P_3	3	6

- **Suppose the scheduling time now is 10**
The Gantt Chart for the schedule is:



- Waiting time for $P_1 = 9$; $P_2 = 30$; $P_3 = 31$
- Average waiting time: $(9 + 30 + 31)/3 = 70/3$
- **Convoy effect** [护航效果]: short process behind long process; or short process has to wait the long process to finish.

FCFS Drawbacks

- A process that does not perform any I/O will monopolize the processor (**Convoy Effect**).
- Favors CPU-bound processes:
 - I/O-bound processes have to wait until CPU-bound process completes.
 - They may have to wait even when their I/O are completed (poor device utilization).
 - We could have kept the I/O devices busy by giving a bit more priority to I/O bound processes.


Shortest Job First (**SJF**) [Optimal]

- Selection function: the process with the shortest expected CPU burst time.
- Decision mode: Non-preemptive.
 - There is a variant which supports “Preemptive”, called **SRJF** (Shortest-Remaining Job First)
- also called Shortest Time First (**STF**) and Shortest Process Next (**SPN**).
- I/O bound processes will be picked first.
- SJF is **optimal** – gives minimum average waiting time for a given set of processes.

Example of SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

- SJF (non-preemptive)



- Average waiting time = $(0 + 6 + 3 + 7)/4 = 4$

Dynamics of Shortest-Job-First (SJF)

- Associate with each process the length of its next CPU burst.
 - Use these lengths to schedule the process with the shortest time.
- **Non-preemptive** – once CPU given to the process it cannot be preempted until completes its CPU burst.
- We need to somehow estimate the required processing time (CPU burst time) for each process.

Determining length of next CPU Burst

- Can only estimate the length.
- Can be done by using the length of previous CPU bursts, using exponential averaging:

1. t_n = actual length of n^{th} CPU burst

2. τ_{n+1} = predicted value for the next CPU burst


3. $\alpha, 0 \leq \alpha \leq 1$

4. Define: $\tau_{n+1} = \alpha t_n + (1 - \alpha) \tau_n.$

Examples of Exponential Averaging

- How to set α in $\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$?
- $\alpha = 0$
 - $\tau_{n+1} = \tau_n$,
 - Recent history does not count.
- $\alpha = 1$
 - $\tau_{n+1} = t_n$,
 - Only the actual last CPU burst counts.
- Let's be balanced: $\alpha = 0.5$
 - See example in next slide.

Prediction of the length of the next CPU Burst



Shortest Job First Drawbacks

- Possibility of **starvation** for longer processes as long as there is a steady supply of shorter processes.
- Lack of preemption is not suited in a time sharing environment:
 - CPU bound process gets lower priority (as it should) but a process doing no I/O could still monopolize the CPU if he is the first one to enter the system.
- SJF implicitly incorporates priorities: shortest jobs are given preferences.

Shortest-Remaining-Job-First (SRJF)

- Idea
 - Associate with each process the length of its next/remaining CPU burst.
 - Use these lengths to schedule the process with the shortest time.
- **Preemptive** – if a new process arrives with CPU burst length less than remaining time of current executing process, preempt.
- Called Shortest-Remaining-Job-First (SRJF).

Example of Preemptive SJF \rightarrow SRTF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
----------------	---------------------	-------------------


P_1	0.0	7
-------	-----	---

P_2	2.0	4
-------	-----	---

P_3	4.0	1
-------	-----	---

P_4	5.0	4
-------	-----	---

- SJF (preemptive)



- Average waiting time = $(9 + 1 + 0 + 2)/4 = 3$

Shortest Remaining Time Scheduling

- **Preemptive version of SJF/SPN**
- Arriving processes with shorter burst preempt a running process
- Very large variance of response times
 - long processes wait even longer than under SPF/SJF/SPN
- Not always optimal
 - Short incoming process can preempt a running process that is near completion
 - Context-switching overhead can become significant

Priority


- A priority number (integer) is assigned for each process
- The CPU is allocated to the process with the highest priority (generally, **smallest integer \equiv highest priority**).
 - Preemptive
 - Non-preemptive
- ➔ SJF is a priority scheduling where priority is the predicted next CPU burst time.
- Problem:
 - **Starvation** – low priority processes may never execute.
- Solution:
 - **Aging** – as time progresses increase the priority of the process.

Example: Priority scheduling (**non-preemptive**)

Process	Arrive Time	Burst Time	Priority
P1	0.0	10	3
P2	4.0	1	1
P3	8.0	2	4
P4	9.0	1	5
P5	11.0	5	2

Priority scheduling (**non-preemptive**)

Arrival time **P1** **P2** **P3** **P4** **P5**
 0 4 8 9 11




Waiting time: $(0+6+8+9+0)/5 = 4.6(\text{ms})$.

Drawback of priority

- The main drawback of priority scheduling:
starvation!
- To avoid:
 - Prevent high priority processes from running indefinitely by changing priorities dynamically:
 - Each process has a base priority
 - increase priorities of waiting processes at each clock tick
 - Preserving the SJF priority
 - increase priorities of i/o bound processes
 - $\text{priority} = 1/f$ (f is the fraction of time-quantum used)


- In fact, FCFS and SJF are two instances of priority scheduling
 - Priority in FCFS is the arrival time
 - Priority in SJF therefore is the execution time of corresponding processes
- The aging strategy can also be used in real life to improve the customer's satisfaction
 - Such as in **bank waiting**
 - We can upgrade the priority of the customers who have been waiting too long time



This strategy is used in the later MLFQ (Multilevel Feedback Queue) model.

Round-Robin (RR) Scheduling


- Round-robin scheduling
 - Based on FIFO
 - Processes run only for a limited amount of time called a **time slice** or **quantum**
 - **Preemptive**
 - Requires the system to maintain several processes in memory to minimize overhead
 - Often used as part of more complex algorithms






Round-Robin (RR)

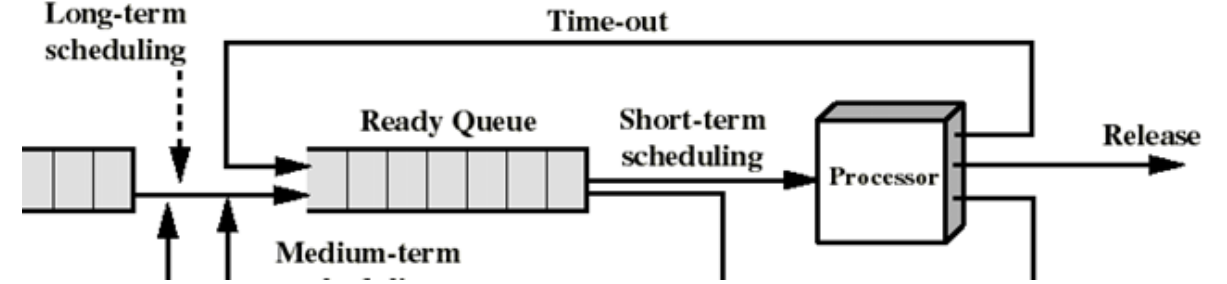
- Selection function: (initially) same as FCFS.
- Decision mode: **preemptive** –
 - a process is allowed to run until the time slice period, called time quantum, is reached.
 - then a clock interrupt occurs and the running process is put at the end of the ready queue.




(a)



(b)




RR Example ($q = 1$)



Picking the Right Quantum

- **Trade-off/Balancing:**
 - Short quantum: great response/interactivity but high overhead
 - Hopefully not too high if the dispatcher is fast enough
 - Long quantum: poor response/interactivity, but low overhead
 - With a very long time quantum, RR Scheduling becomes FCFS Scheduling
- If context-switching time is 10% of time quantum, then the CPU spends >10% of its time doing context switches
- In practice, %CPU time spent on switching is very low
 - time quantum: 10ms to 100ms
 - context-switching time: 10 μ s

Picking the Right Quantum



RR Discussion

- Advantages
 - Jobs get fair share of CPU
 - Shortest jobs finish relatively quickly
- Disadvantages
 - Poor average waiting time with similar job lengths
 - Example: 10 jobs each requiring 10 time slices
 - RR: All complete after about 100 time slices
 - FCFS performs better!
 - Performance depends on length of time-slice
 - If time-slice too short, pay overhead of context switch
 - If time-slice too long, degenerate to FCFS

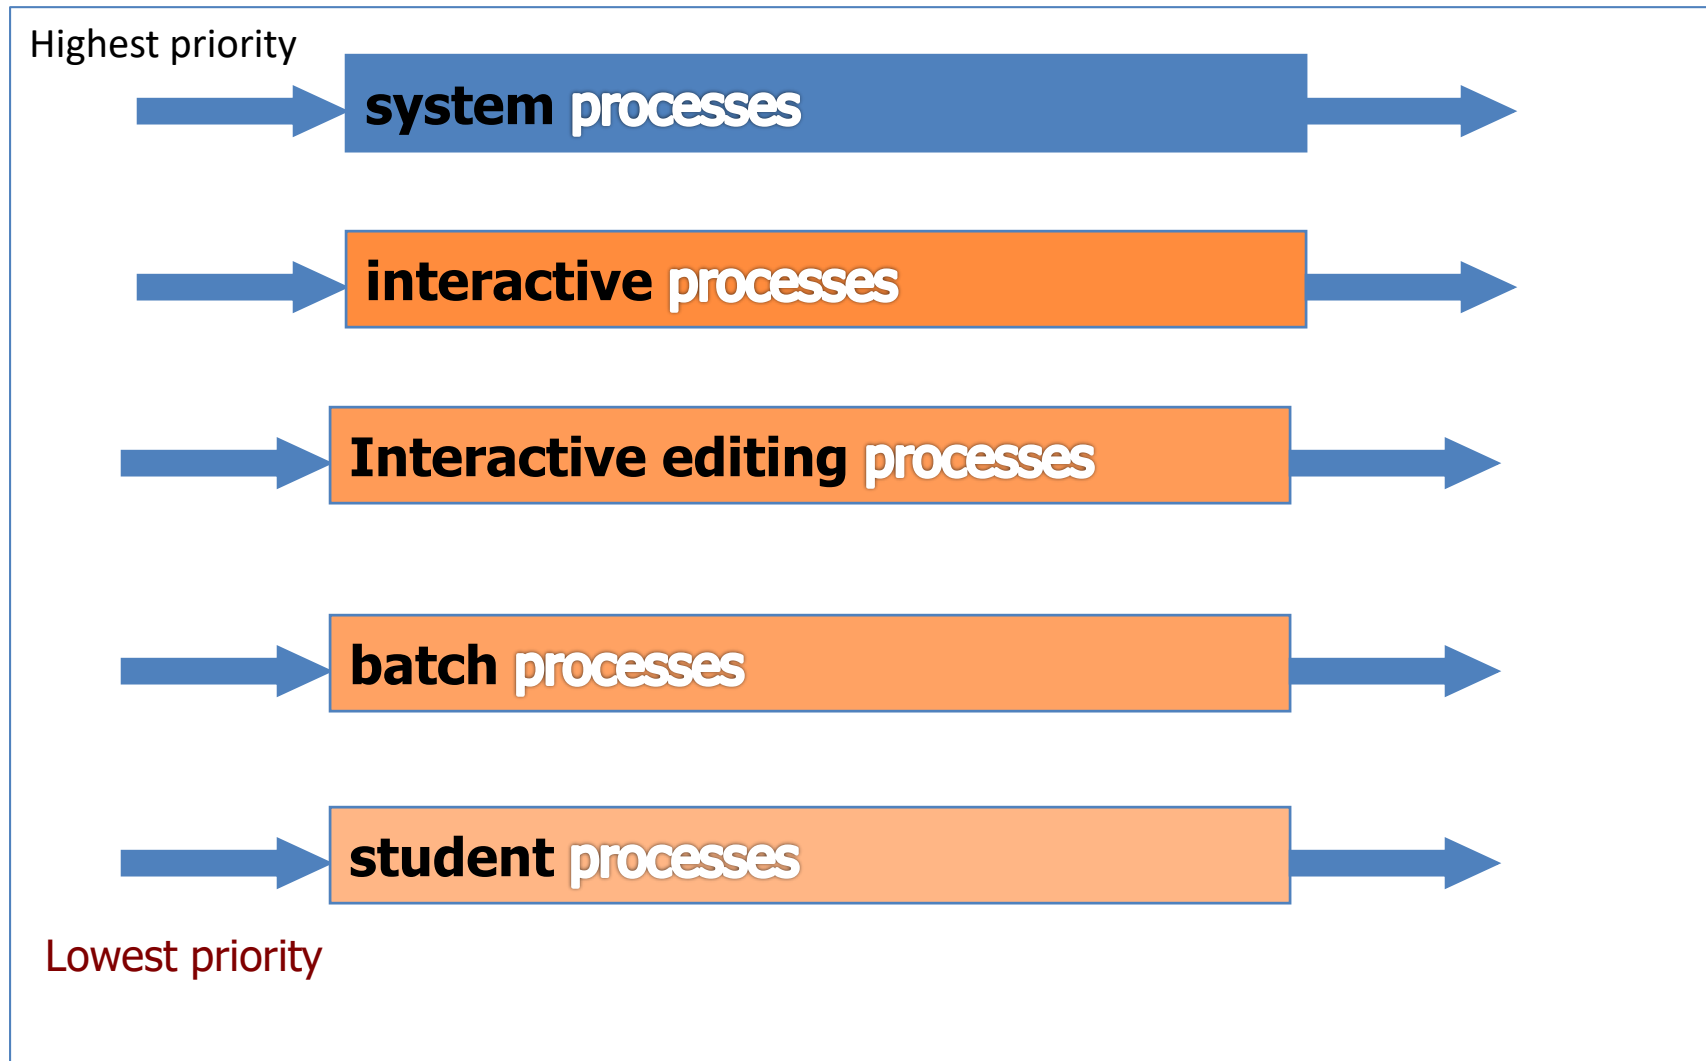
Multilevel Queue

- The RR Scheduling scheme treats all processes equally
- In practice, one often wants to classify processes in groups, e.g., based on externally-defined process priorities
 - Simple idea: use one ready queue per class of processes
 - e.g., if we support 10 priorities, we maintain 10 ready queues
- **Scheduling within queues**
 - Each queue has its own scheduling policy
 - e.g., High-priority could be RR, Low-priority could be FCFS
- **Scheduling between the queues**
 - Typically preemptive priority scheduling
 - A process can run only if all higher-priority queues are empty
 - Or time-slicing among queues
 - e.g., 80% to Queue #1 and 20% to Queue #2

Example of Multilevel Queue

- Ready queue is partitioned into separate queues:
 - foreground (**interactive**)
 - background (**batch**)
- Each queue has its own scheduling algorithm,
 - foreground – **RR**
 - background – **FCFS**
- Scheduling must be done between the queues.
 - Fixed priority scheduling; (i.e., serve all from foreground then from background). Possibility of starvation.
 - Time slice – each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e.,
 - 80% to foreground in RR
 - 20% to background in FCFS

Multi-Level Queue Scheduling



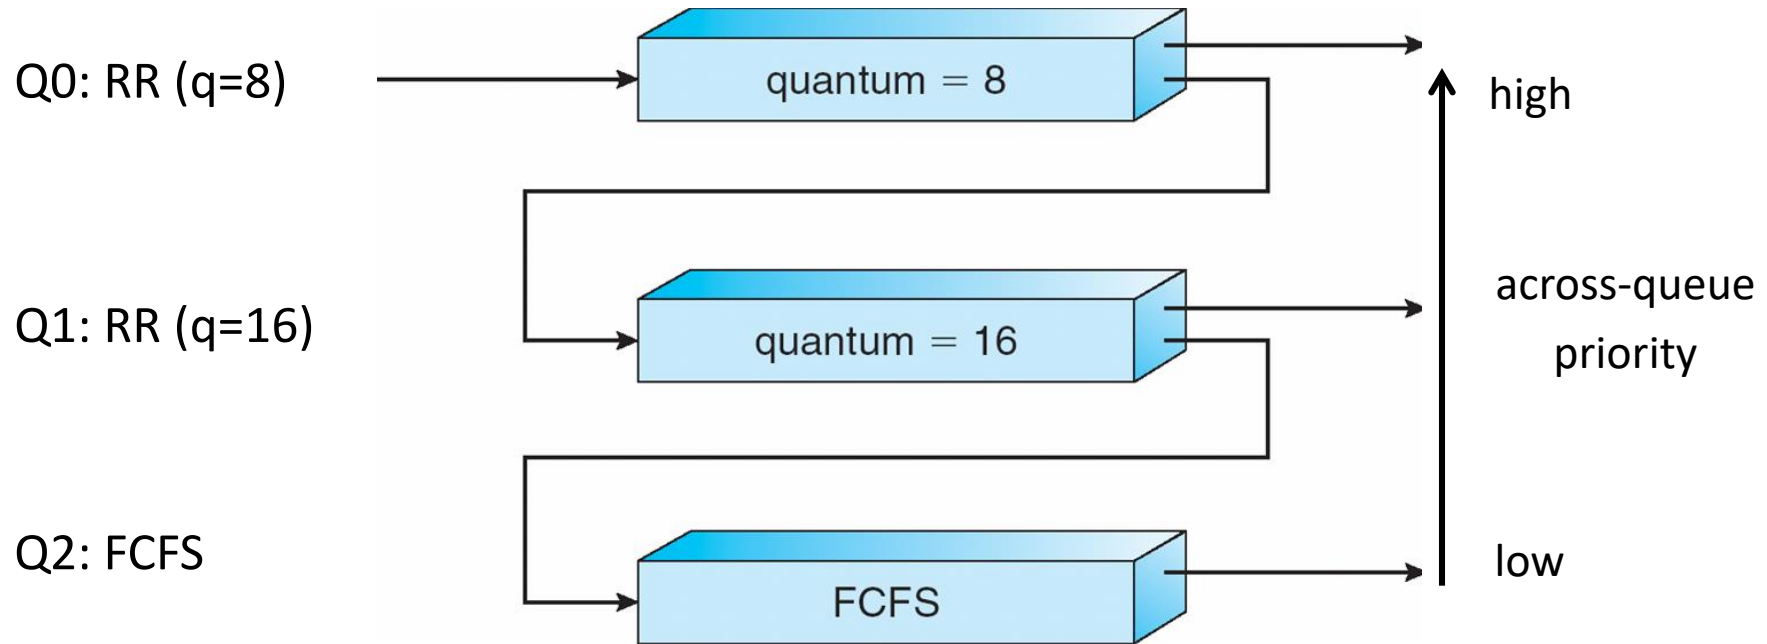
Multilevel Feedback Queue

- Processes can move among the queues
 - If queues are defined on internal process characteristics, it makes sense to move a process whose characteristics have changed
 - e.g., based on CPU burst length
 - It's also a good way to implement priority aging

Example of Multilevel Feedback Queue


- Three queues:
 - Q_0 – RR with time quantum 8 milliseconds
 - Q_1 – RR time quantum 16 milliseconds
 - Q_2 – FCFS
- Scheduling
 - A new job enters queue Q_0 which is served FCFS. When it gains CPU, job receives 8 milliseconds. If it does not finish in 8 milliseconds, job is moved to queue Q_1 .
 - At Q_1 , job is again served FCFS and receives 16 additional milliseconds. If it still does not complete, it is preempted and moved to queue Q_2 .

Example of Multilevel Feedback Queue



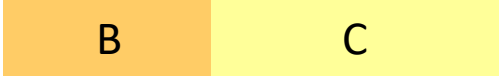

Multilevel Feedback Queues

- Priority 0 (time slice = 1):
- Priority 1 (time slice = 2):
- Priority 2 (time slice = 4):



Multilevel Feedback Queues



time = 1

- Priority 0 (time slice = 1): 
- Priority 1 (time slice = 2): 
- Priority 2 (time slice = 4):



Multilevel Feedback Queues

time = 2

- Priority 0 (time slice = 1): 
- Priority 1 (time slice = 2): 
- Priority 2 (time slice = 4):



Multilevel Feedback Queues

time = 3

- Priority 0 (time slice = 1):
- Priority 1 (time slice = 2):
- Priority 2 (time slice = 4):



Multilevel Feedback Queues

time = 3

- Priority 0 (time slice = 1):
- Priority 1 (time slice = 2):
- Priority 2 (time slice = 4):



suppose process A is blocked on an I/O



Multilevel Feedback Queues

time = 3

- Priority 0 (time slice = 1):
- Priority 1 (time slice = 2):
- Priority 2 (time slice = 4):

A

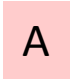
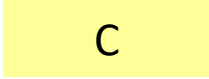
B C

suppose process A is blocked on an I/O



Multilevel Feedback Queues

time = 5

- Priority 0 (time slice = 1): 
- Priority 1 (time slice = 2): 
- Priority 2 (time slice = 4):

suppose process A is returned from an I/O



Multilevel Feedback Queues

time = 6

- Priority 0 (time slice = 1):
- Priority 1 (time slice = 2):
- Priority 2 (time slice = 4):


c



Multilevel Feedback Queues

time = 8


- Priority 0 (time slice = 1):
- Priority 1 (time slice = 2):
- Priority 2 (time slice = 4): c



Multilevel Feedback Queues

time = 9


- Priority 0 (time slice = 1):
- Priority 1 (time slice = 2):
- Priority 2 (time slice = 4):



Lottery Scheduling

- Give every job some number of **lottery tickets**.
- On each time slice, randomly pick a winning ticket.
- On average, CPU time is proportional to the number of tickets given to each job.
- Assign tickets by giving the most to short running jobs, and fewer to long running jobs (approximating SJF).
To avoid starvation, every job gets at least one ticket.
- Degrades gracefully as load changes. Adding or deleting a job affects all jobs proportionately, independent of the number of tickets a job has.

Lottery Scheduling Example



9


$P1=6$

$P2=9$

1	4	7	10	13
2	5	8	11	14
3	6	9	12	15

Schedule P2

Lottery Scheduling Example



3

$P1=6$

1

4

2

5

3

6

$P2=9$

7

10

13

8

11

14


9

12

15

Schedule P1

Lottery Scheduling Example



- As $t \rightarrow \infty$, processes will get their share (unless they were blocked a lot)
- Problem with Lottery scheduling: Only probabilistic guarantee
- What does the scheduler have to do
 - When a new process arrives?
 - When a process terminates?

Schedule P2

Highest Response Ratio Next (HRRN)

- aims to minimize **Response ratio** (T_q/T_s) for each process
 - Response ratio (**RR** later) $T_q/T_s = (T_w + T_s)/T_s$
 - where T_w = waiting time
 T_s = expected service time
- can approximate an *a priori* measure :
 - expected service time must be estimated again
 - waiting time measured as time progresses

- Basic Concepts
- Scheduling Criteria & Metrics
- Different Scheduling Algorithms
- Algorithm Evaluation

Algorithms Comparison

- Which one is best?
- The answer depends on:
 - on the system workload (extremely variable).
 - hardware support for the dispatcher.
 - relative weighting of performance criteria (response time, CPU utilization, throughput...).
 - The evaluation method used (each has its limitations...).
- Hence the answer depends on too many factors to give any...

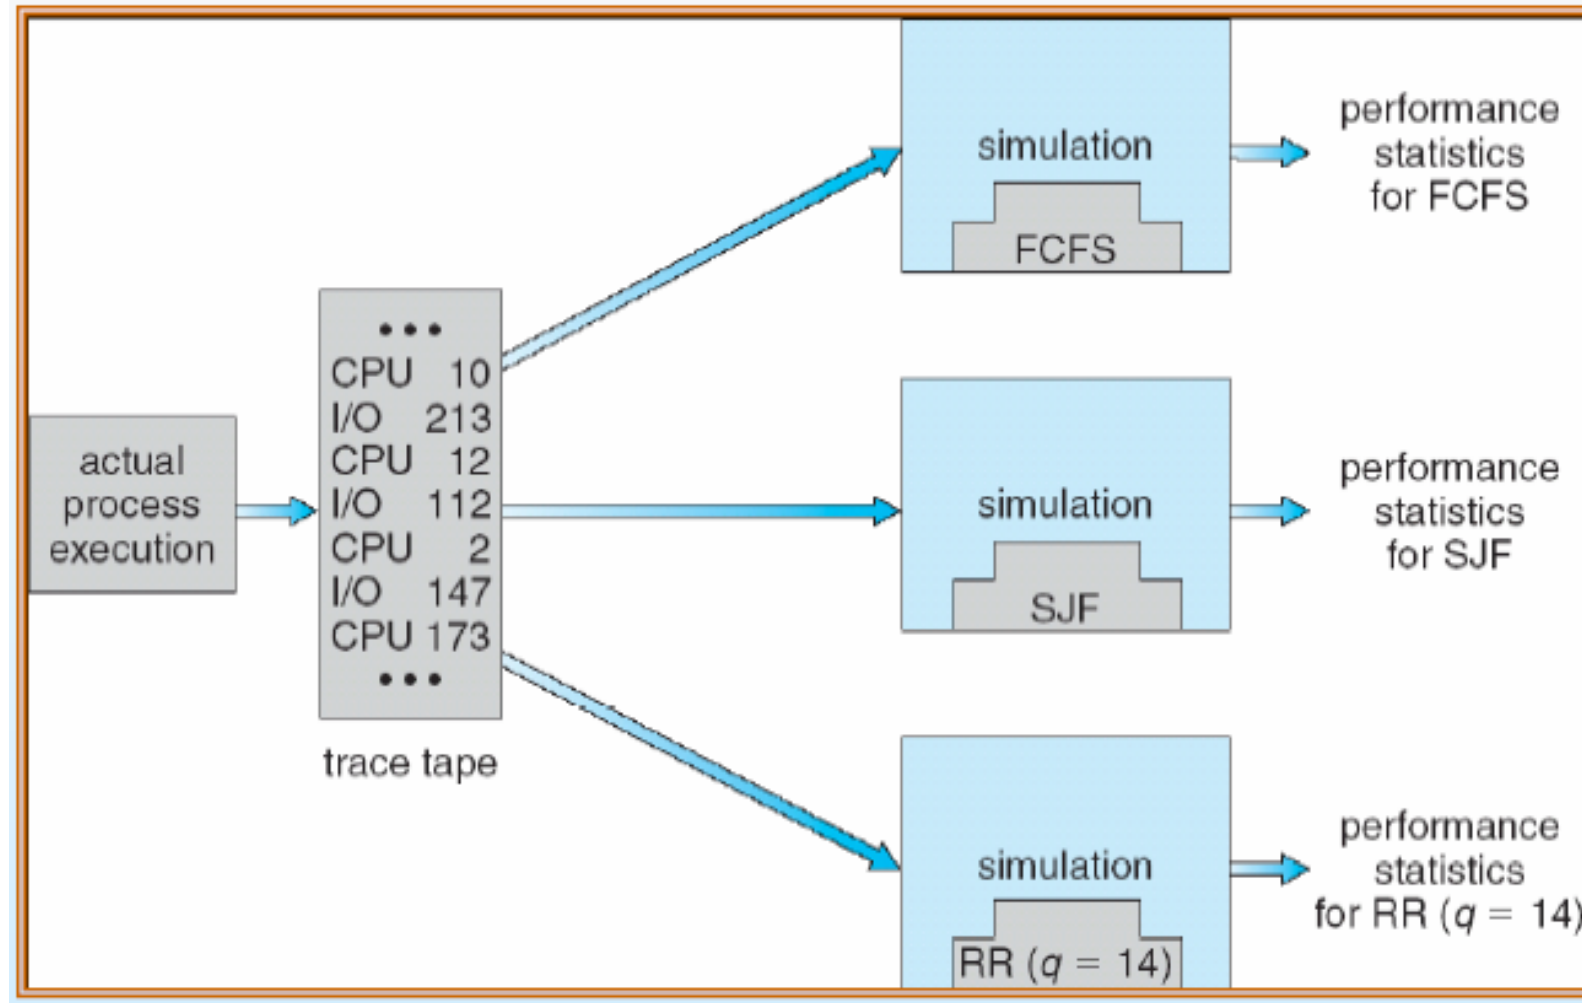
Deterministic modeling

- Deterministic modeling [确定模型法] takes a particular predetermined workload and defines the performance of each algorithm for that workload.
 - To describe scheduling algorithms and provide examples,
- Deterministic model is simple and fast. It gives the exact numbers, allowing us to compare the algorithms.
- However, it requires exact numbers for input, and its answers apply only to these cases.

Simulations

- Simulations [模拟] involve programming a model of the computer system.
 - Software **data structures** represent the major components of the system.
 - The **simulator** has a variable representing a clock; as this variable's value is increased, the simulator modifies the system to reflect the activities of the device, the processes, and the scheduler.
 - As the simulation executes, **statistics** that indicate algorithm performance are gathered and printed.
- Artificial data or trace tapes.

Simulations



Simulation - Queuing models

- Queuing models [排队模型]

- **Queueing-network analysis**

- The computer system is described as a network of servers. Each server has a queue of waiting process. The CPU is a server with its ready queue, as is the I/O system with its ready queue.
 - Knowing arrival rates and service times, we can find utilization, average queue length, and average response time.

- Useful for comparing scheduling algorithms. However, probability distributions are difficult to work with, and some assumptions required.

Power of mathematics –
this is the general way to
find solutions for
problems

- ICQ A [10 pts]
 - Next week
 - 1 hour, close
 - Blank-filling, Computation