

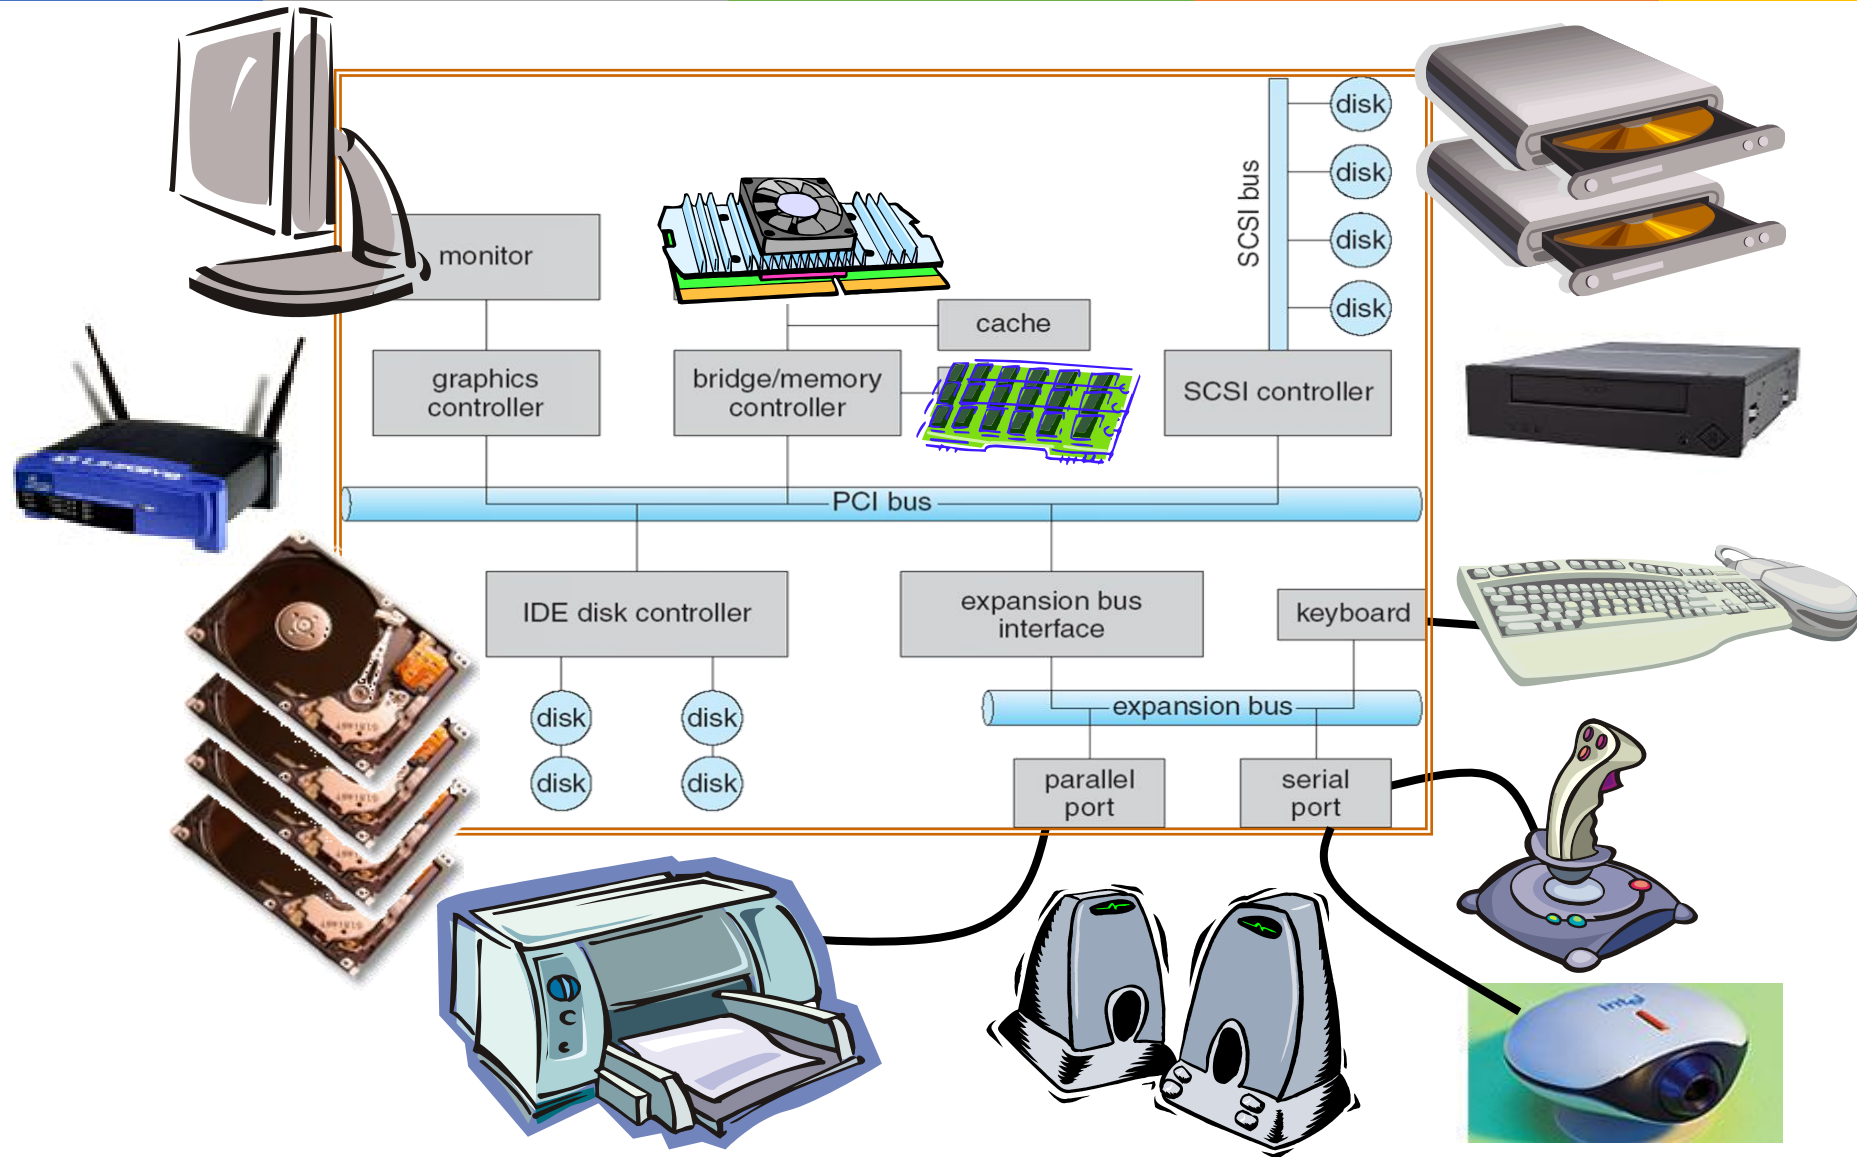


北京交通大学

I/O System



Modern I/O Systems



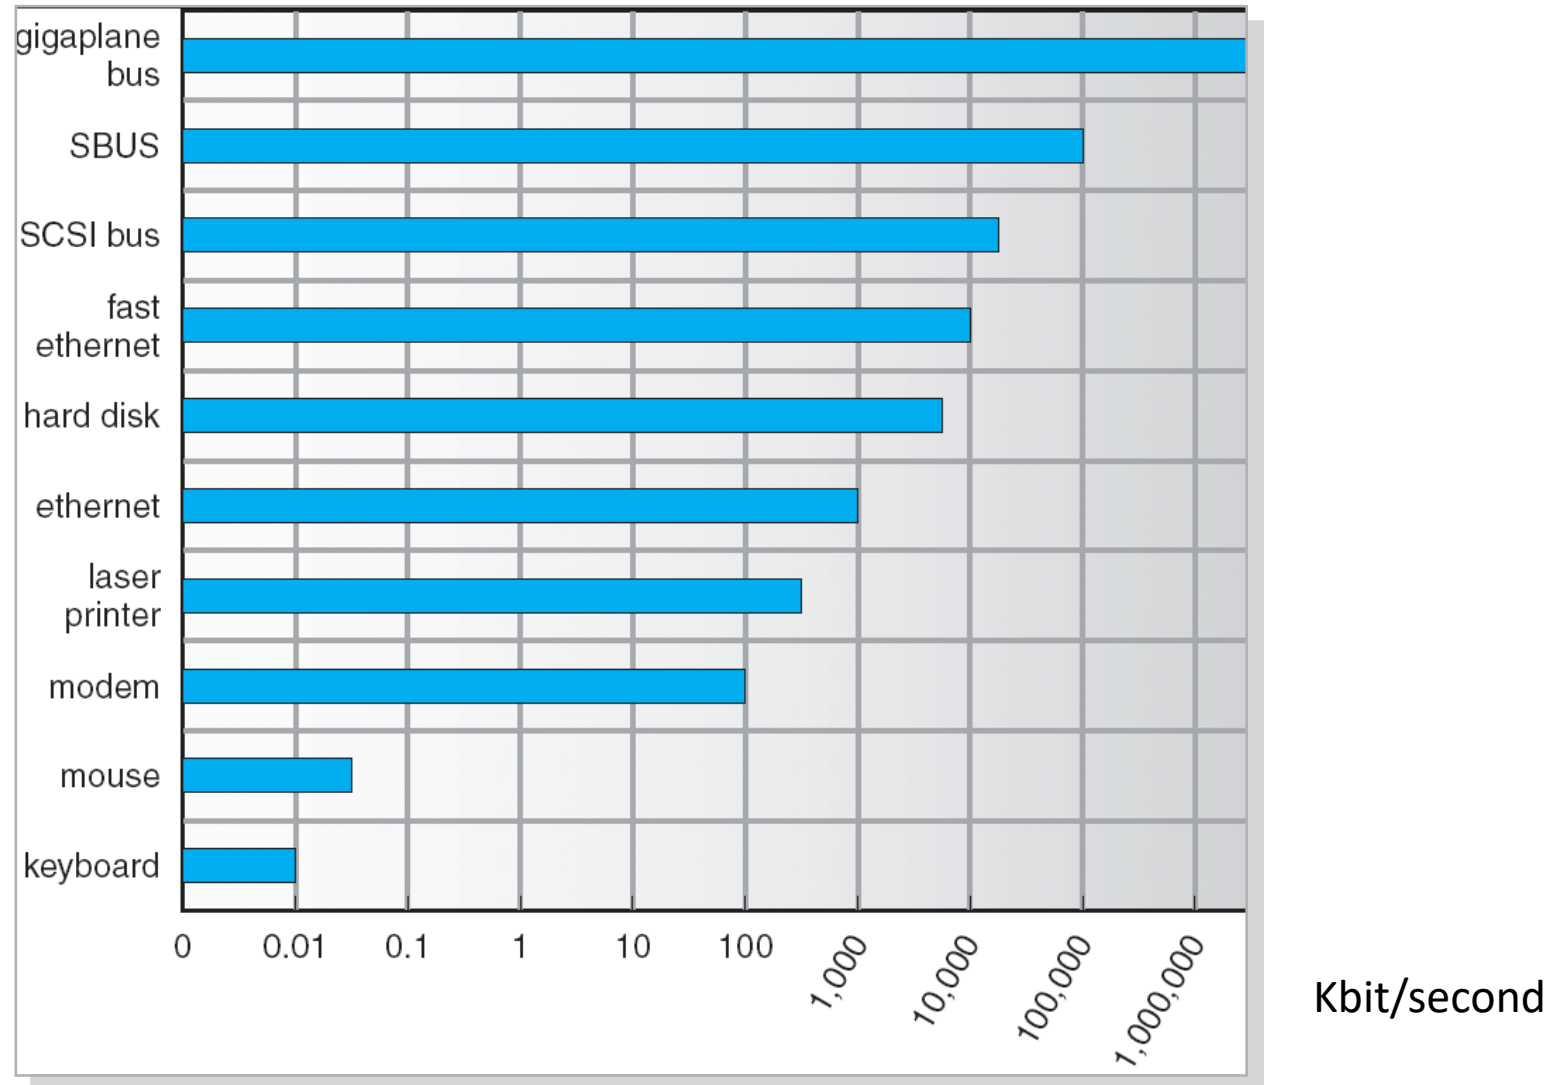
The Requirements of I/O

- So far in this course:
 - We have learned how to manage CPU, memory.....
- What about I/O?
 - Without I/O, computers are useless
 - But...
 - Thousands of devices, each slightly different
 - Devices unreliable: media failures and transmission errors
 - Devices unpredictable and/or slow

Three Main Problems

- Thousands of devices, each slightly different
 - How can we standardize the interfaces to these devices?
- Devices unreliable: media failures and transmission errors
 - How can we make them reliable???
- Devices unpredictable and/or slow
 - How can we manage them if we don't know what they will do or how they will perform?

Example: Device-Transfer Rates



I/O Devices

- Devices vary in many dimensions
 - **Character-stream** or **block**
 - **Sequential** or **random-access**
 - **Sharable** or **dedicated**
 - **Speed** of **operation**
 - **Read-write, read only, or write only**

I/O Devices

- **Block Devices:** e.g. disk drives, tape drives, DVD-ROM
 - Access blocks of data
 - Commands include `open()`, `read()`, `write()`, `seek()`
- **Character Devices:** e.g. keyboards, mouse, serial ports, some USB devices
 - Single characters at a time
 - Commands include `get()`, `put()`
- **Network Devices:** e.g. Ethernet, Wireless, Bluetooth
 - Different enough from block/character to have own interface

I/O Devices

- Sequential/Random
 - Some devices must be accessed sequentially (*e.g.* tape)
 - Others can be accessed randomly (*e.g.* Disk, CD, etc.)

Characteristics of I/O Devices

aspect	variation	example
data-transfer mode	character block	terminal disk
access method	sequential random	modem CD-ROM
transfer schedule	synchronous asynchronous	tape keyboard
sharing	dedicated sharable	tape keyboard
device speed	latency seek time transfer rate delay between operations	
I/O direction	read only write only read–write	CD-ROM graphics controller disk

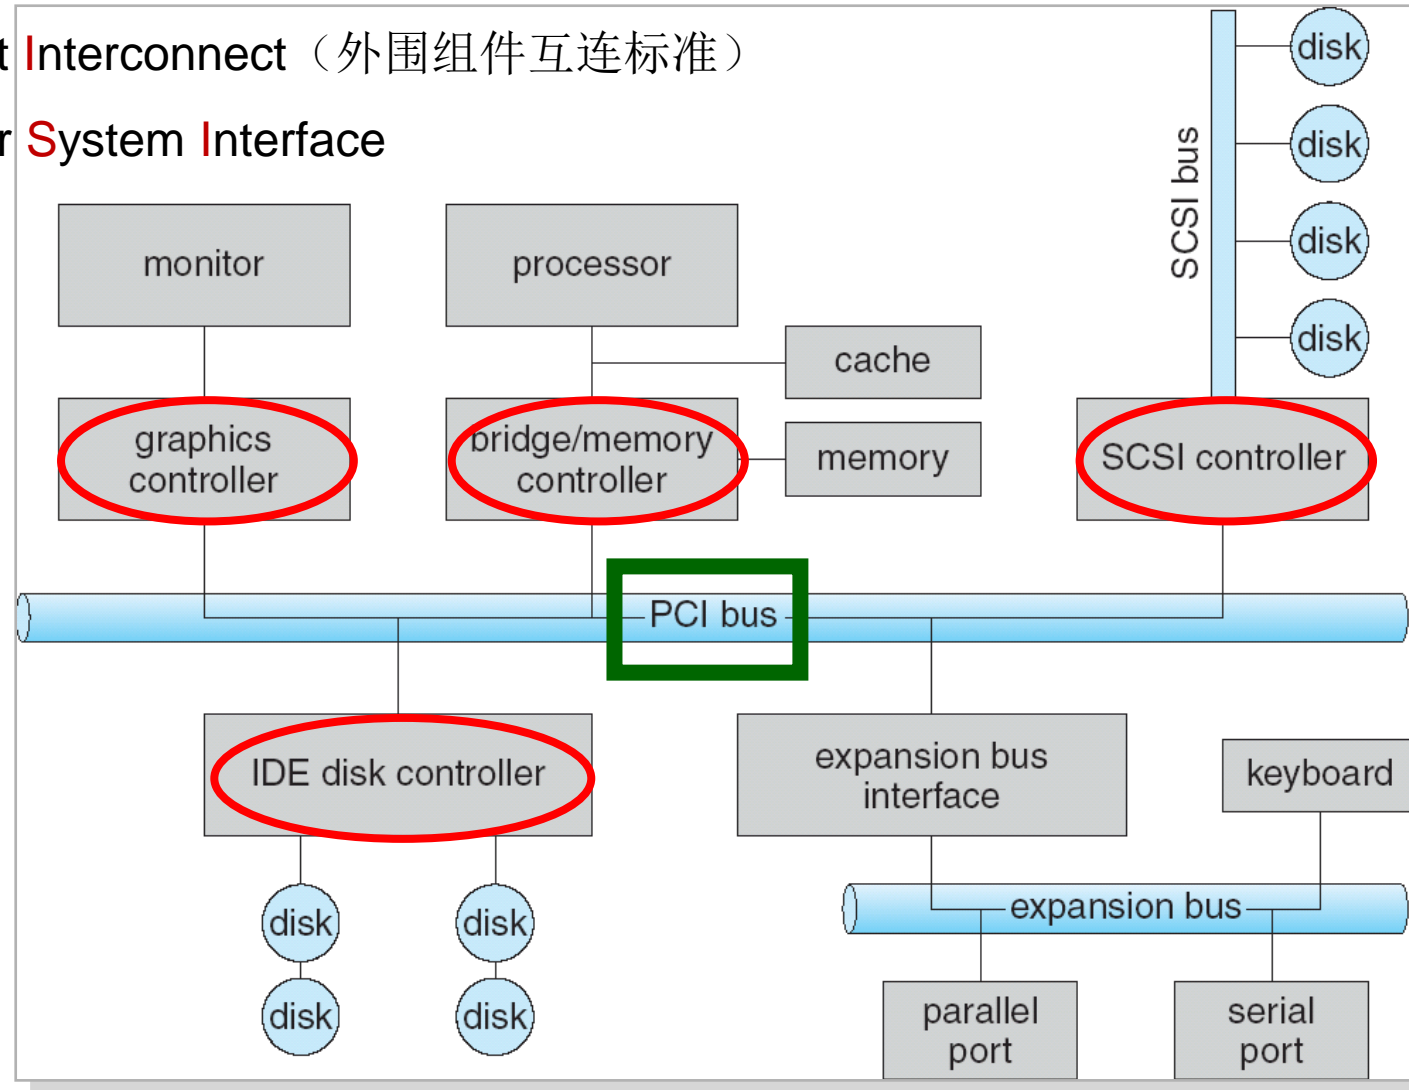
I/O Hardware

- Incredible variety of I/O devices
- Common concepts
 - **Port** -- address
 - **Bus**
 - **Controller (host adapter)**
 - Contains a set of registers that can be read and written
 - May contain memory
- I/O instructions control devices
- Devices have addresses

A Typical PCI Bus Structure

Peripheral Component Interconnect (外围组件互连标准)

SCSI: Small Computer System Interface

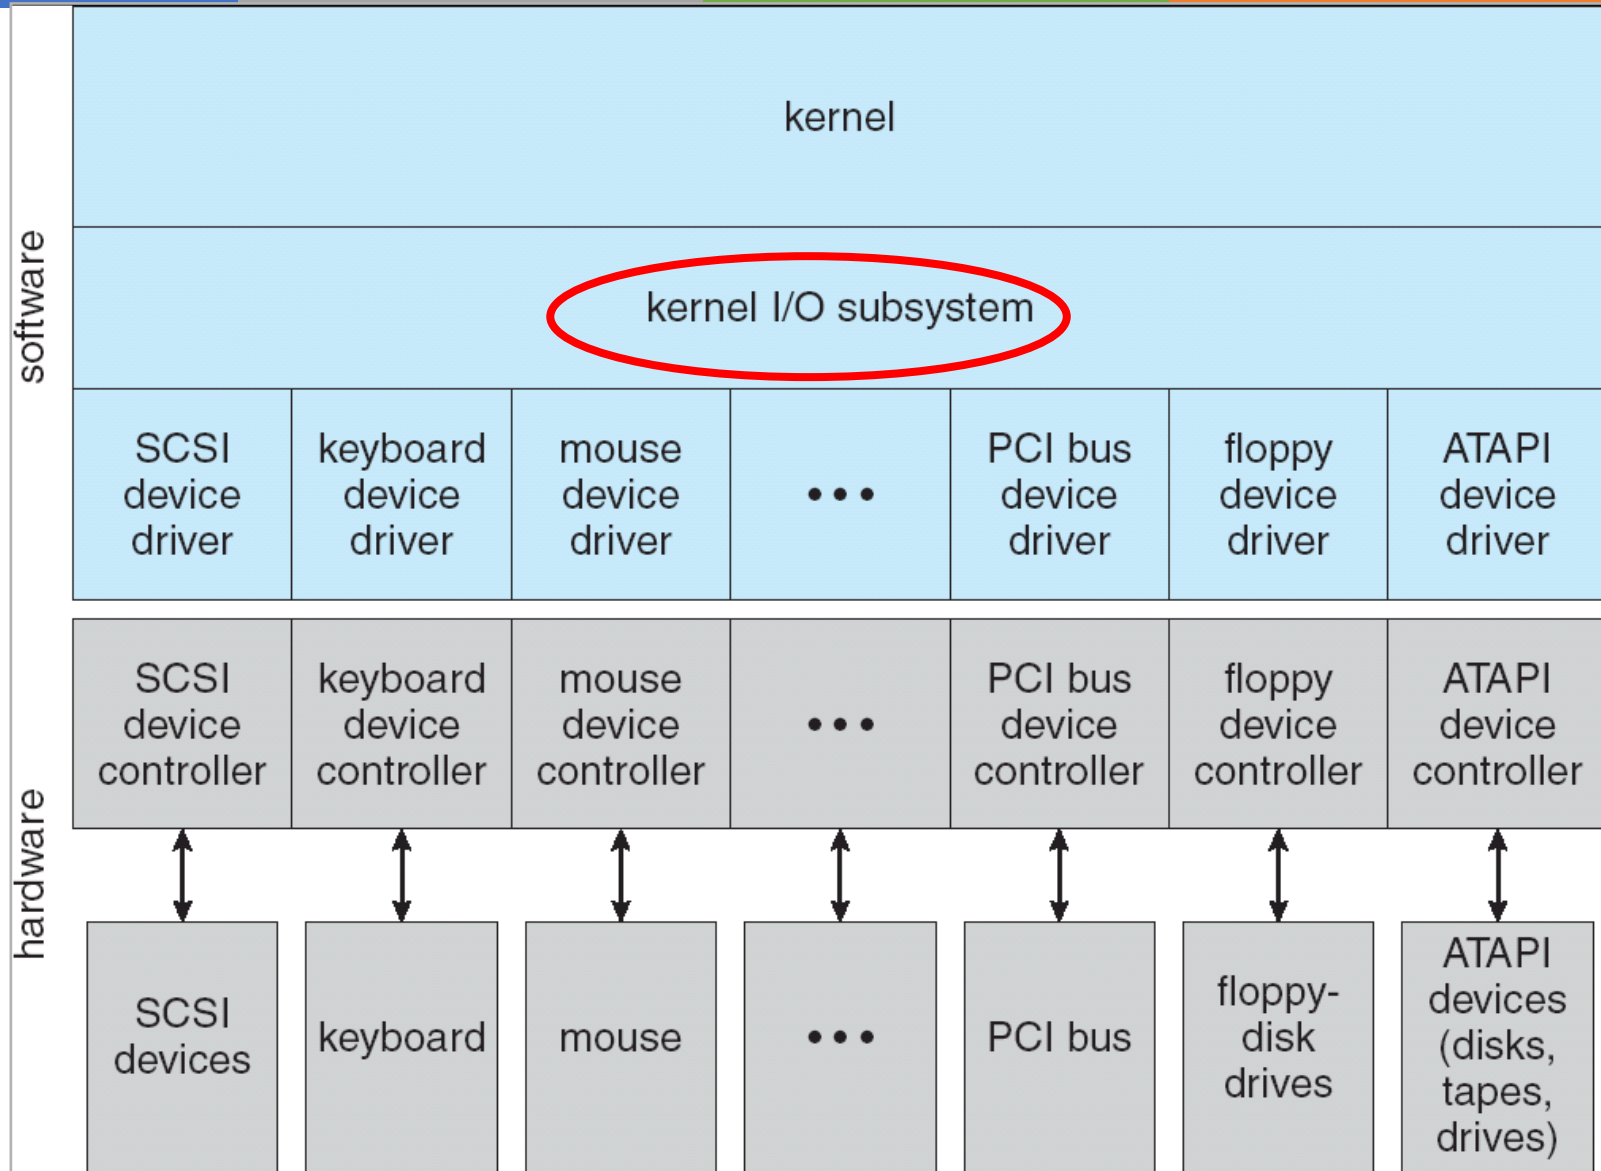


Device I/O Port Locations on PCs (partial)



I/O address range (hexadecimal)	device
000–00F	DMA controller
020–021	interrupt controller
040–043	timer
200–20F	game controller
2F8–2FF	serial port (secondary)
320–32F	hard-disk controller
378–37F	parallel port
3D0–3DF	graphics controller
3F0–3F7	diskette-drive controller
3F8–3FF	serial port (primary)

A Kernel I/O Structure



Device-driver layer
hides differences
among I/O
controllers from
kernel

The Goal of I/O Subsystem

- Provide **Uniform** Interfaces, Despite Wide Range of Different Devices

- This code works on many different devices:

```
int fd = open("/dev/something") ;  
for (int i = 0; i < 10; i++) {  
    fprintf(fd, "Count %d\n", i) ;  
}  
close(fd) ;
```

- Why? Because code that controls devices ("device driver") implements **standard** interface.

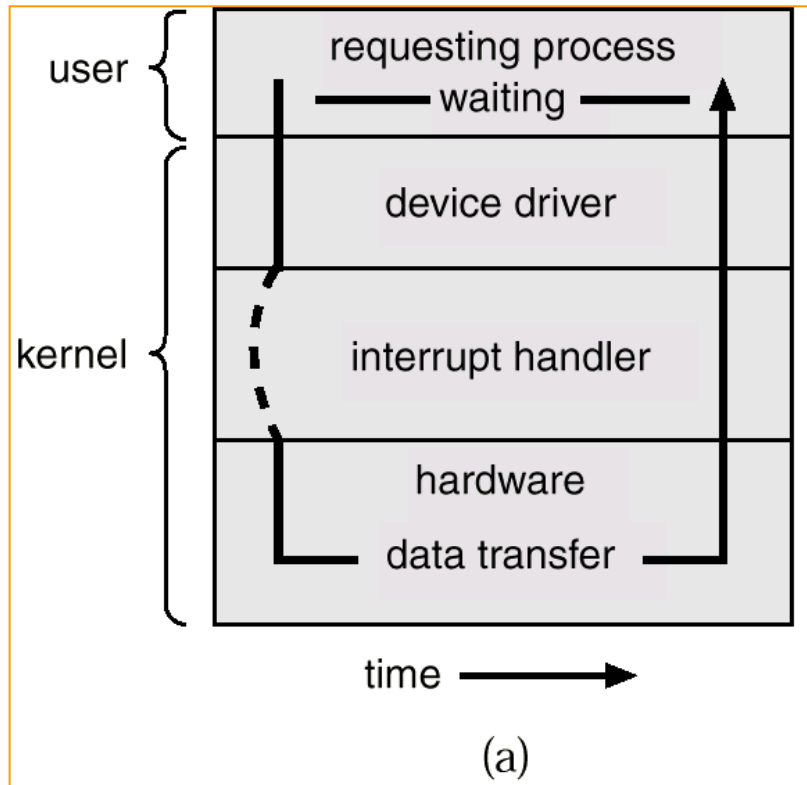
Three Main Problems

- Incredible variety of I/O devices.
 - Unpredictable and/or slow.
 - Unreliable.
-
- I/O organization
 - Device management

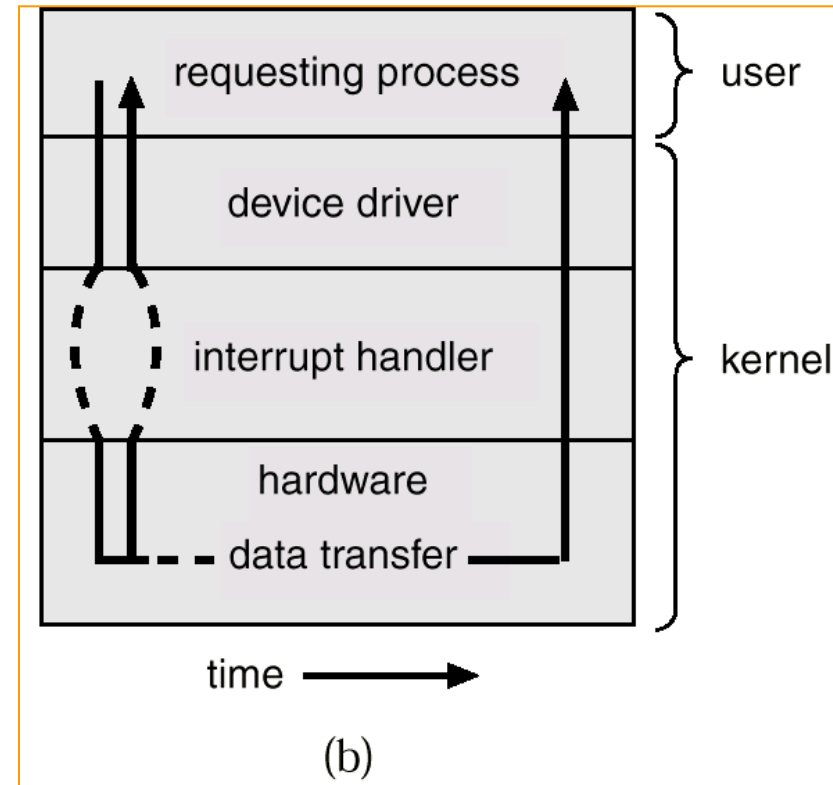


I/O Organization

Two I/O Methods



Synchronous



Asynchronous

Two I/O Methods

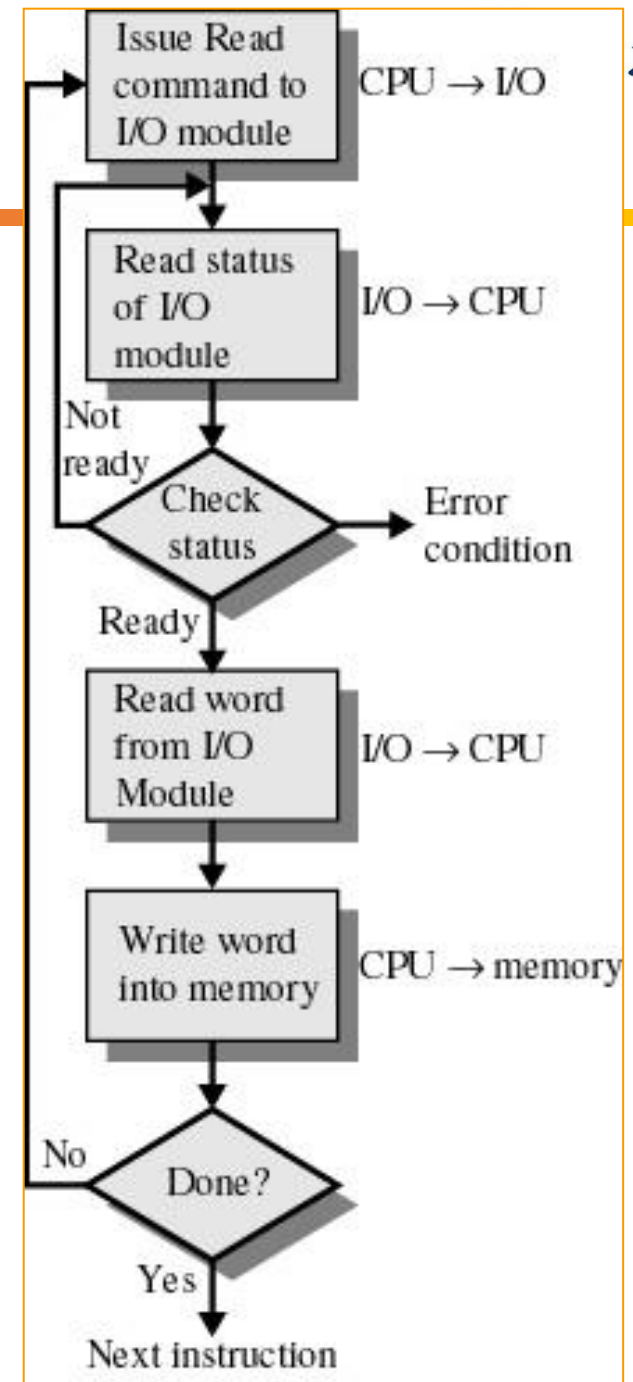
- **Synchronous** - process suspended until I/O completed
 - Easy to use and understand
 - Insufficient for some needs
- **Asynchronous** - process runs while I/O executes
 - Difficult to use
 - I/O subsystem signals process when I/O completed

Two I/O Methods

- **Synchronous: "Wait"**
 - When request data (e.g. `read()` system call), put process to sleep until data is ready
 - When write data (e.g. `write()` system call), put process to sleep until device is ready for data
- **Asynchronous Interface: "Tell Me Later"**
 - When request data, take pointer to user's buffer, return immediately; later kernel fills buffer and notifies user
 - When send data, take pointer to user's buffer, return immediately; later kernel takes data and notifies user

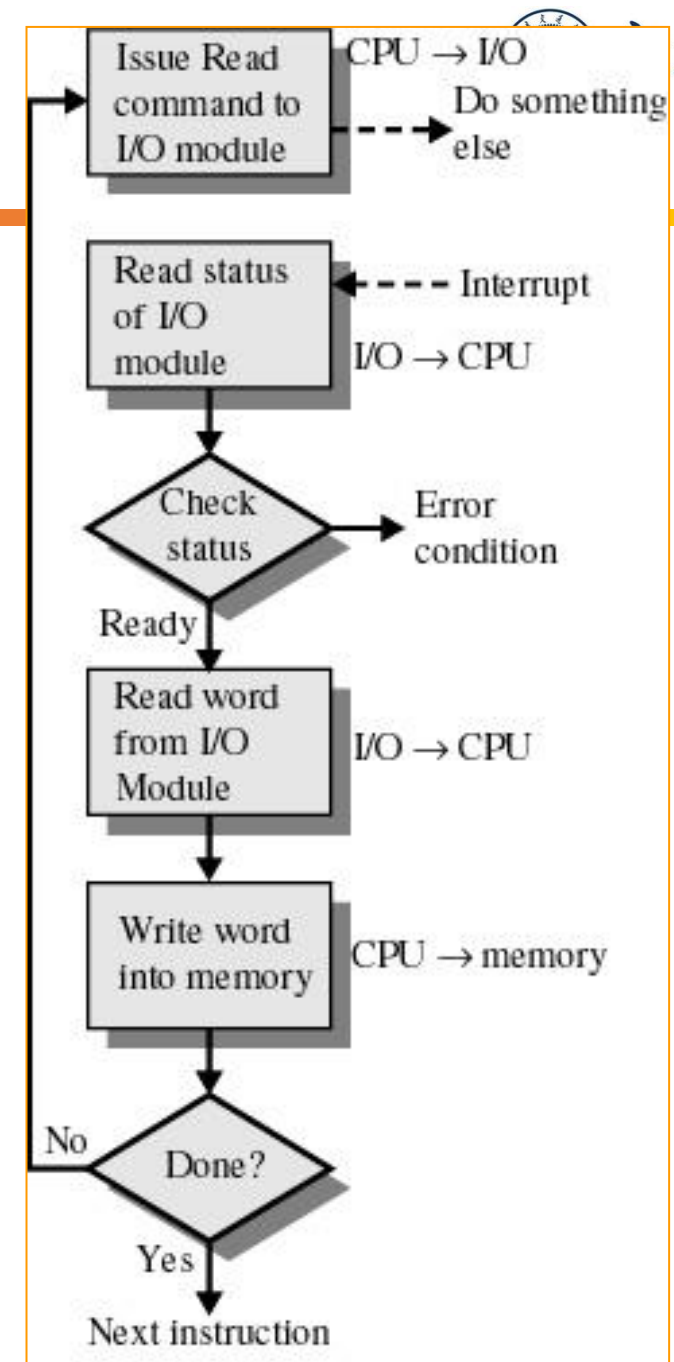
Polling

- Determines state of device
 - Ready or Busy or Error
- The host repeatedly reads the busy bit until that bit becomes clear

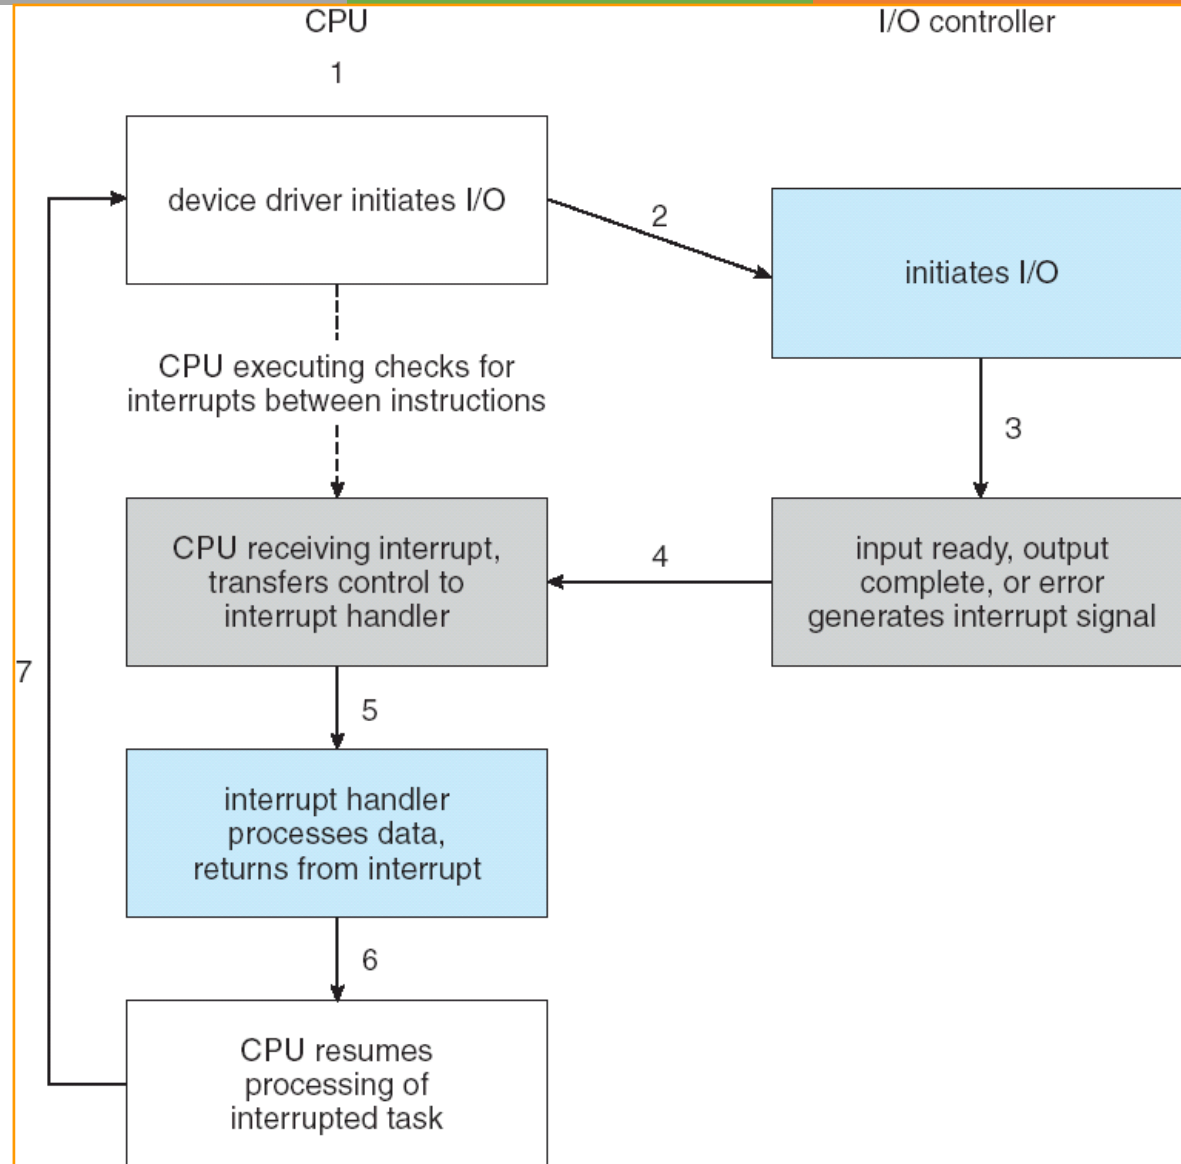


Interrupts

- CPU is responsible for
 - moving chars to/from controller buffer, but
- Interrupt signal informs CPU when I/O operation completes



Interrupt-Driven I/O Cycle

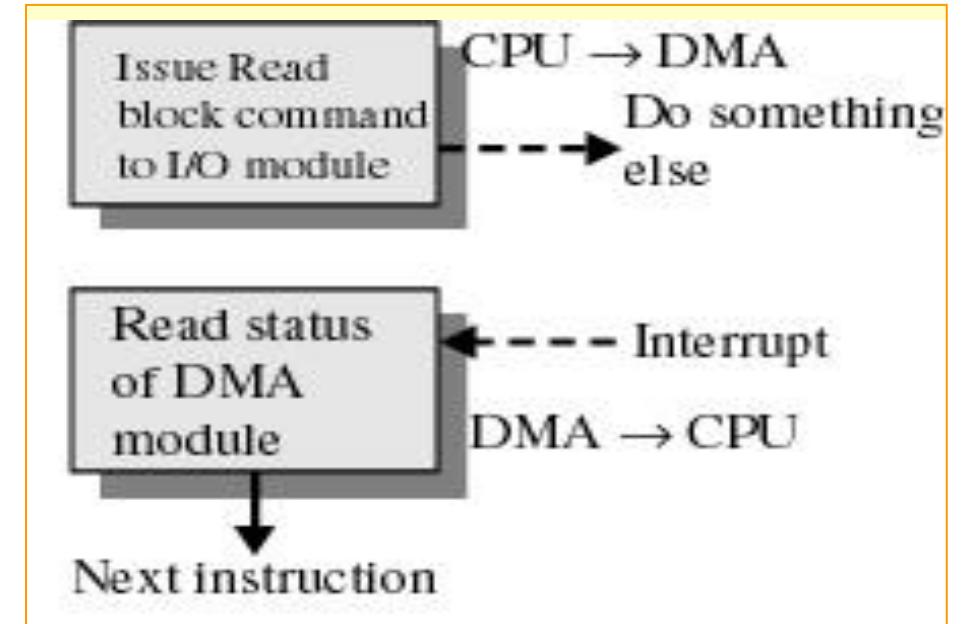


Direct Memory Access

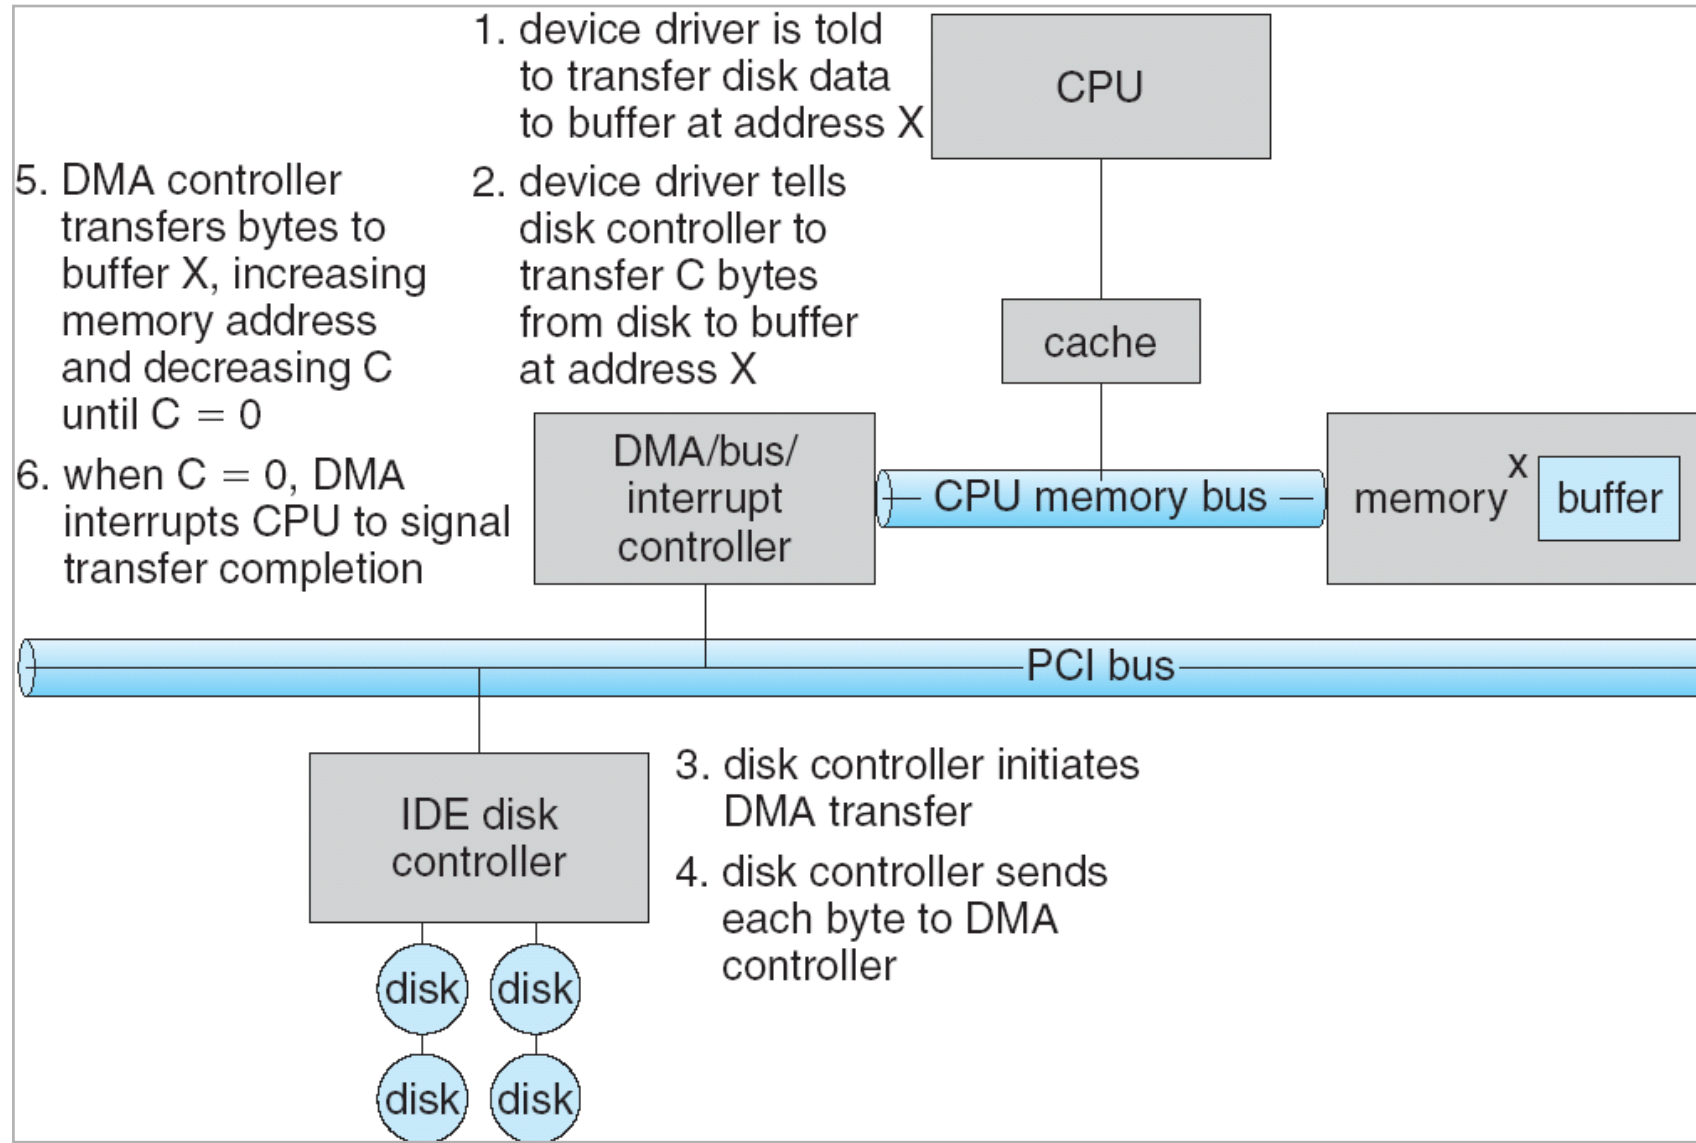
- Used to avoid **programmed I/O** for large data movement
- Requires **DMA** controller
- Bypasses CPU to transfer data directly between I/O device and memory

DMA

- CPU only initiates operation
- DMA controller transfers data directly to/from main memory
- Interrupt when transfer completed



Six Step Process to Perform DMA Transfer





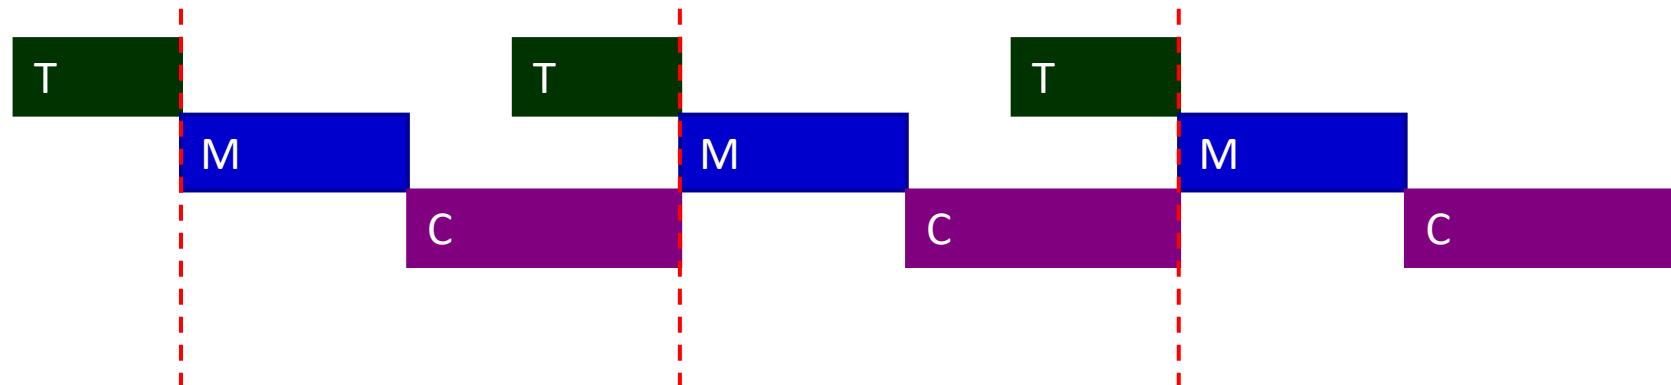
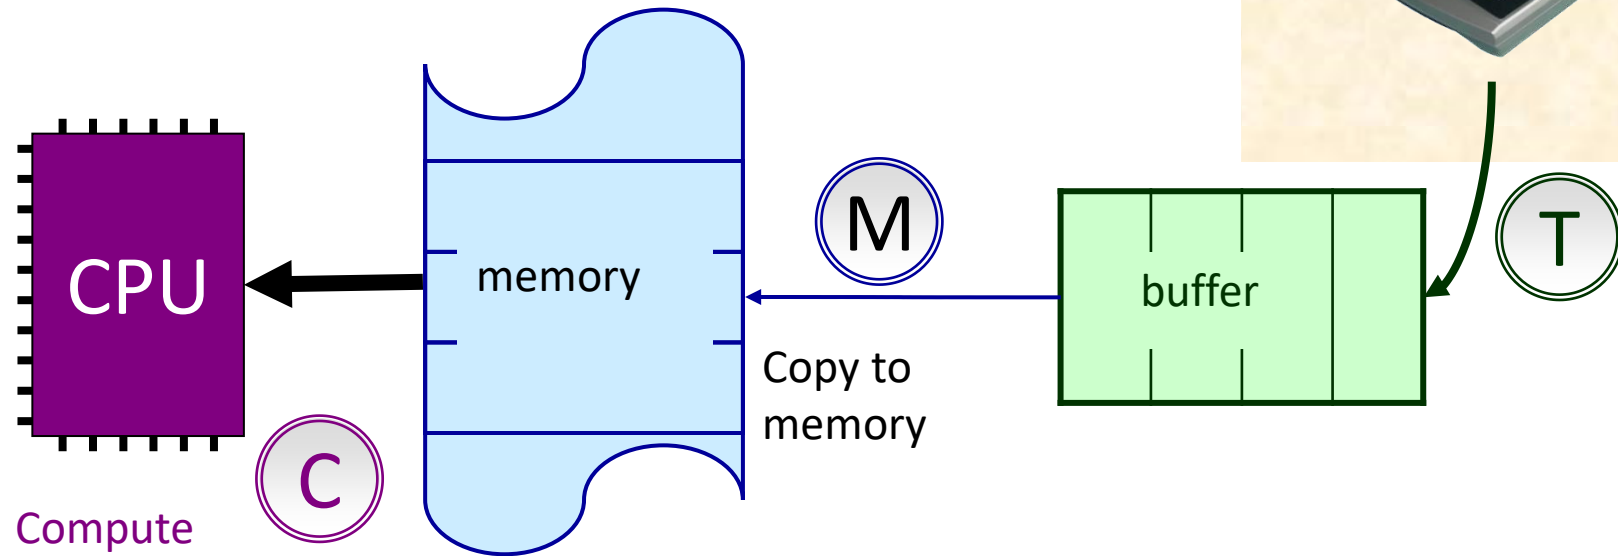
Device Management

Device Management

- Scheduling
- Error handling
- Buffering - store data in memory while transferring between devices
 - To cope with device speed mismatch

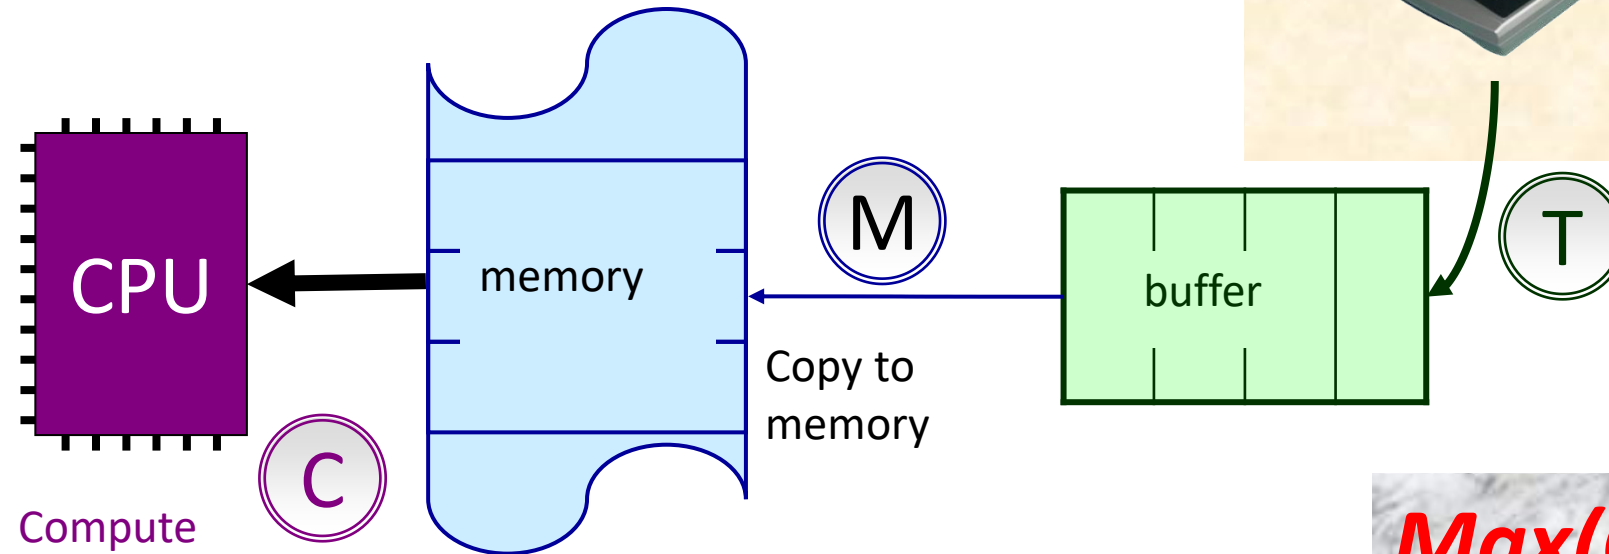
Device Management

Single Buffer

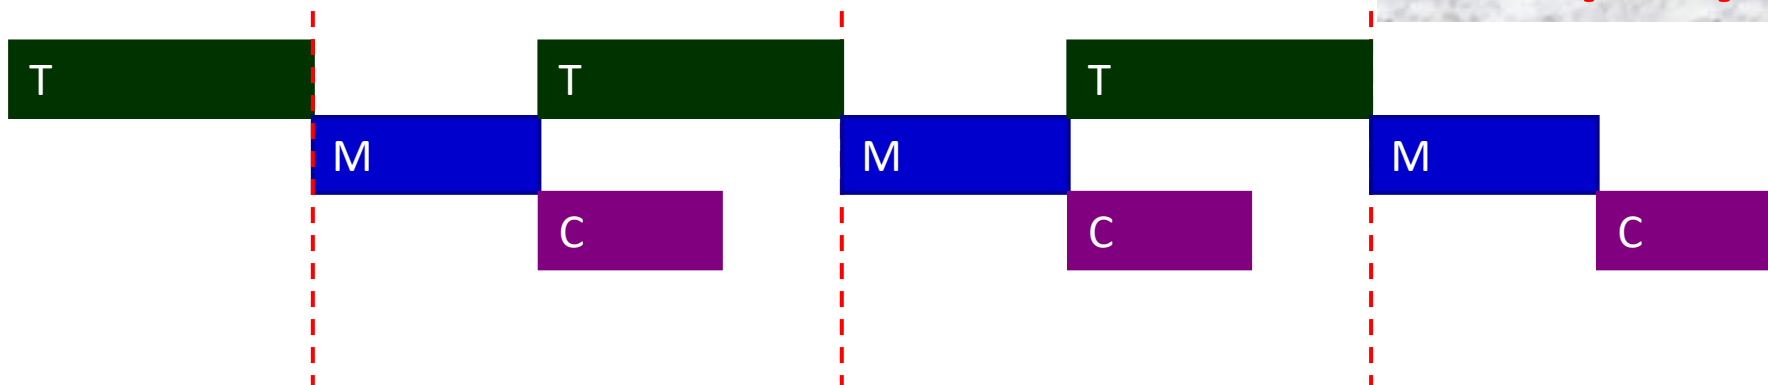


Device Management

Single Buffer



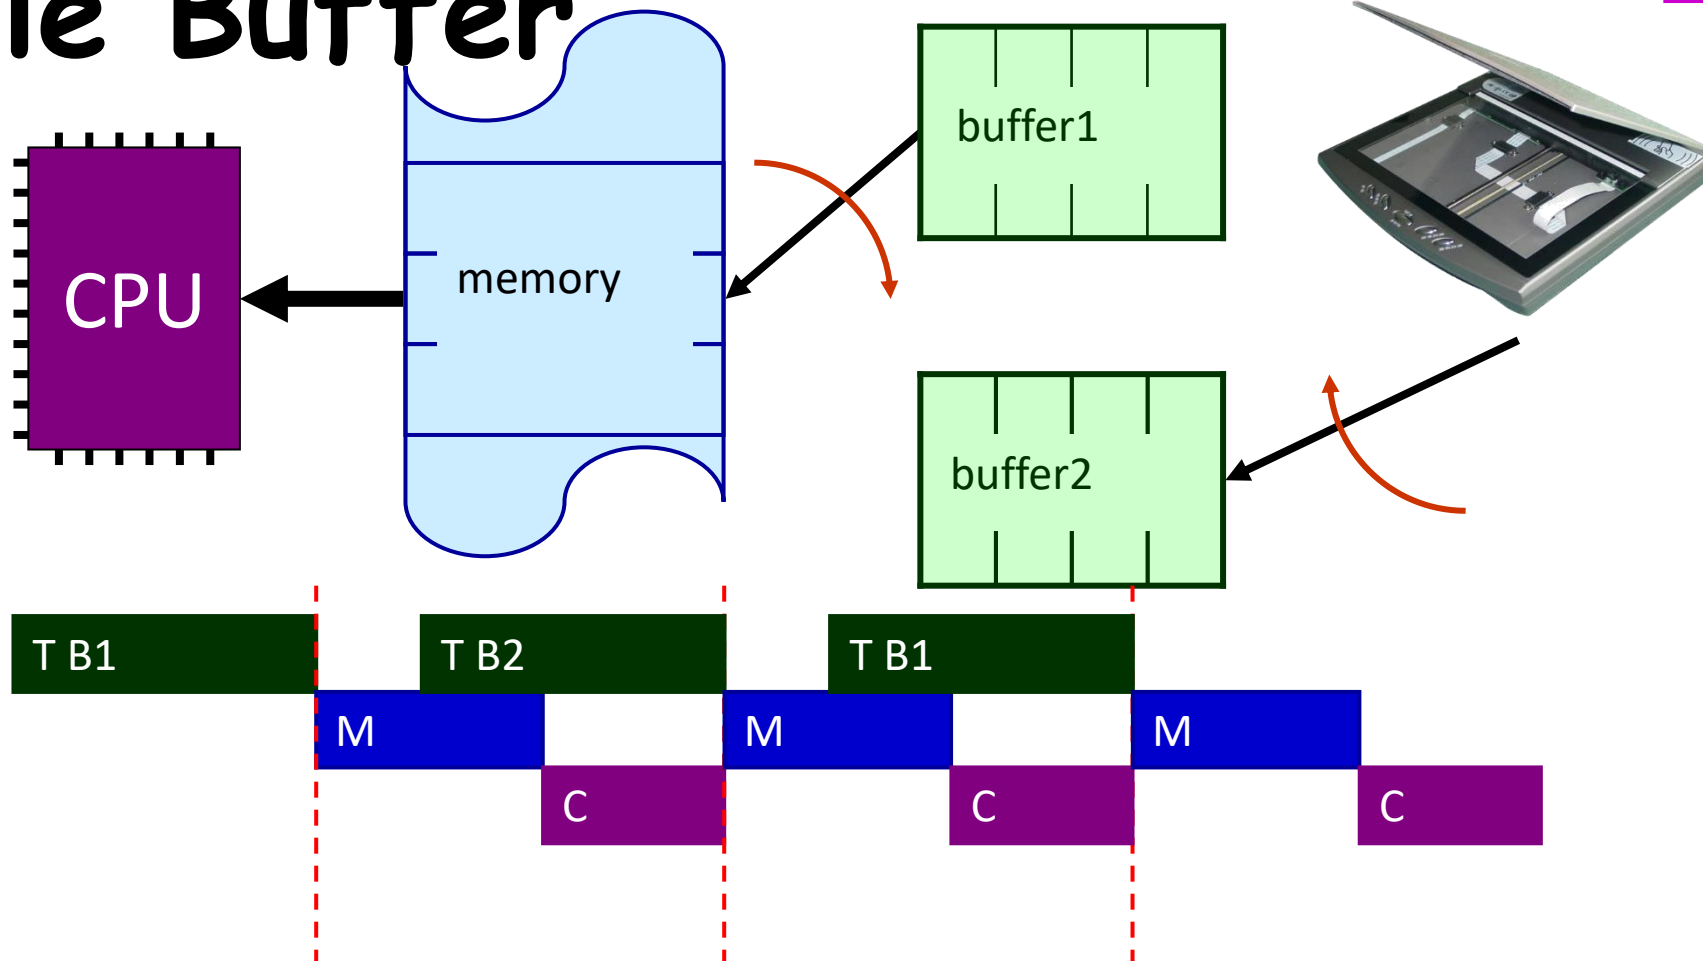
$$Max(C,T)+M$$



Device Management

$T > C$

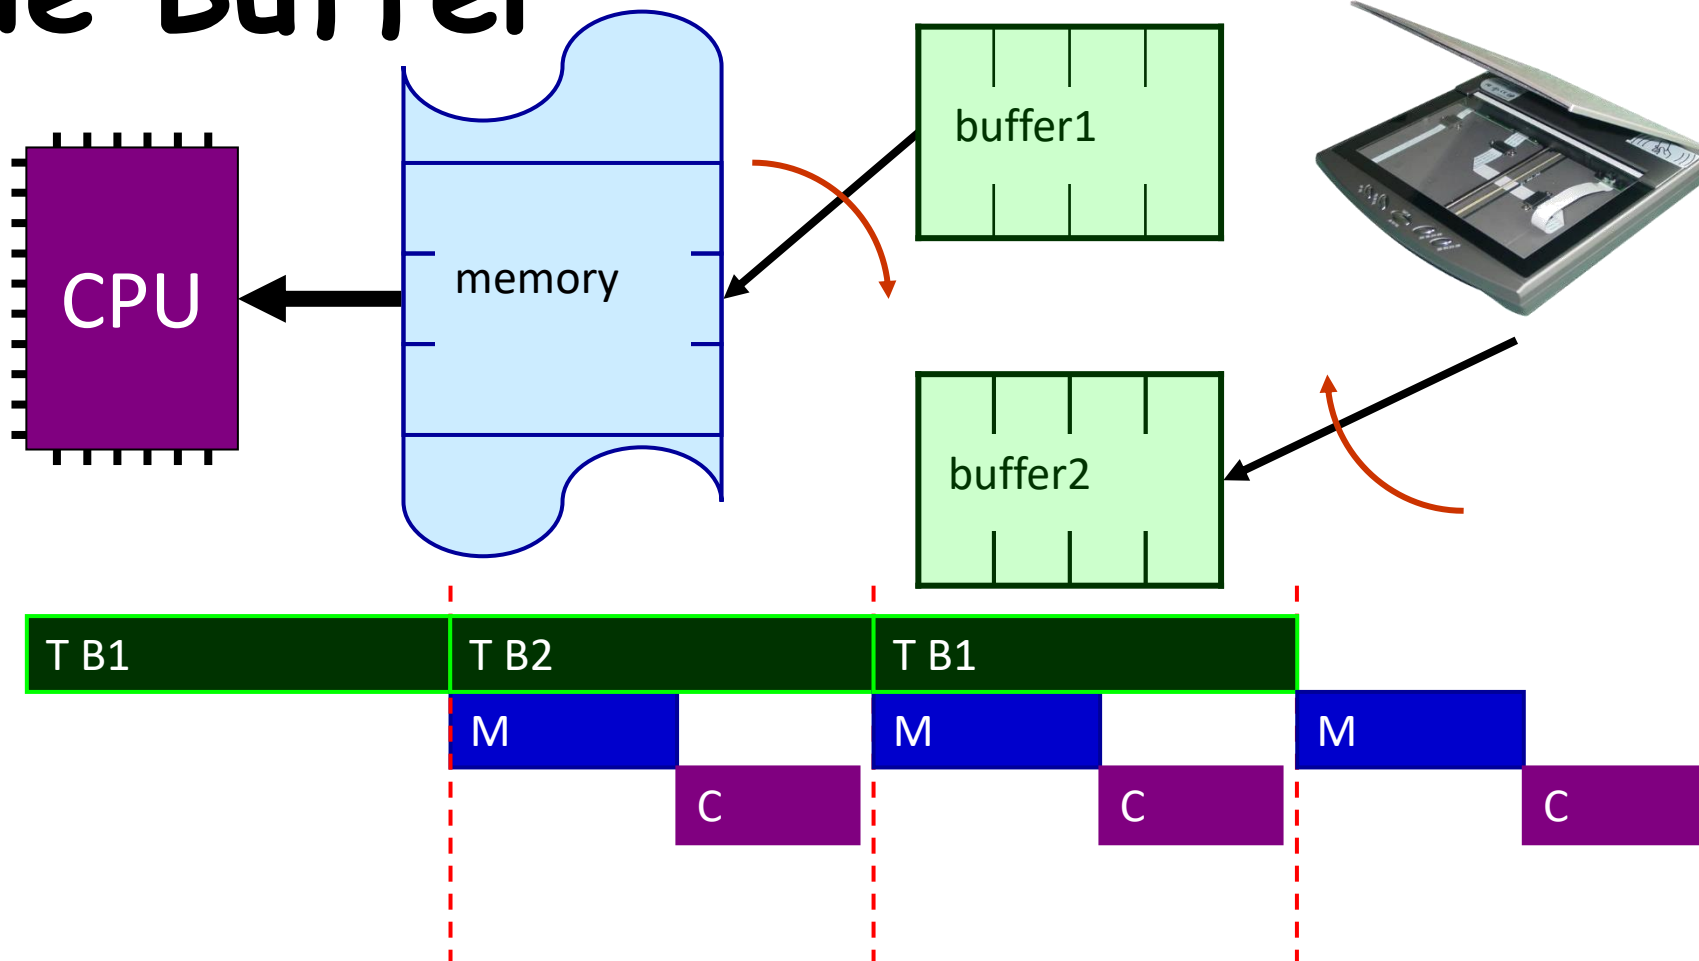
Double Buffer



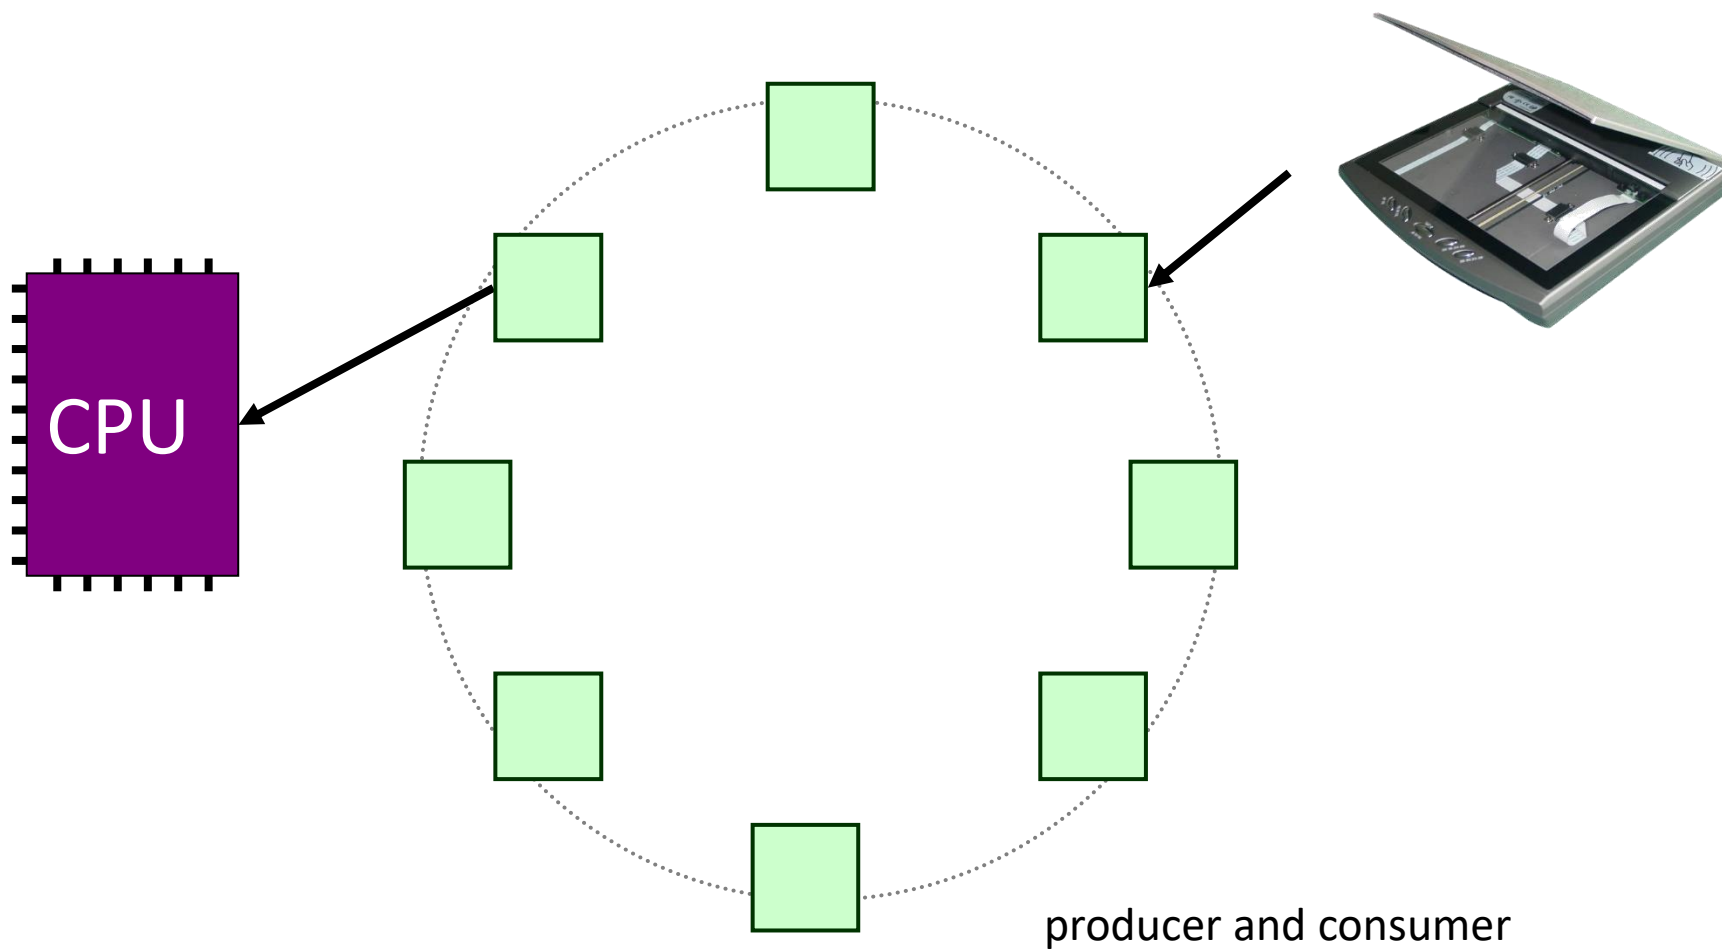
Device Management

Double Buffer

$T > C$



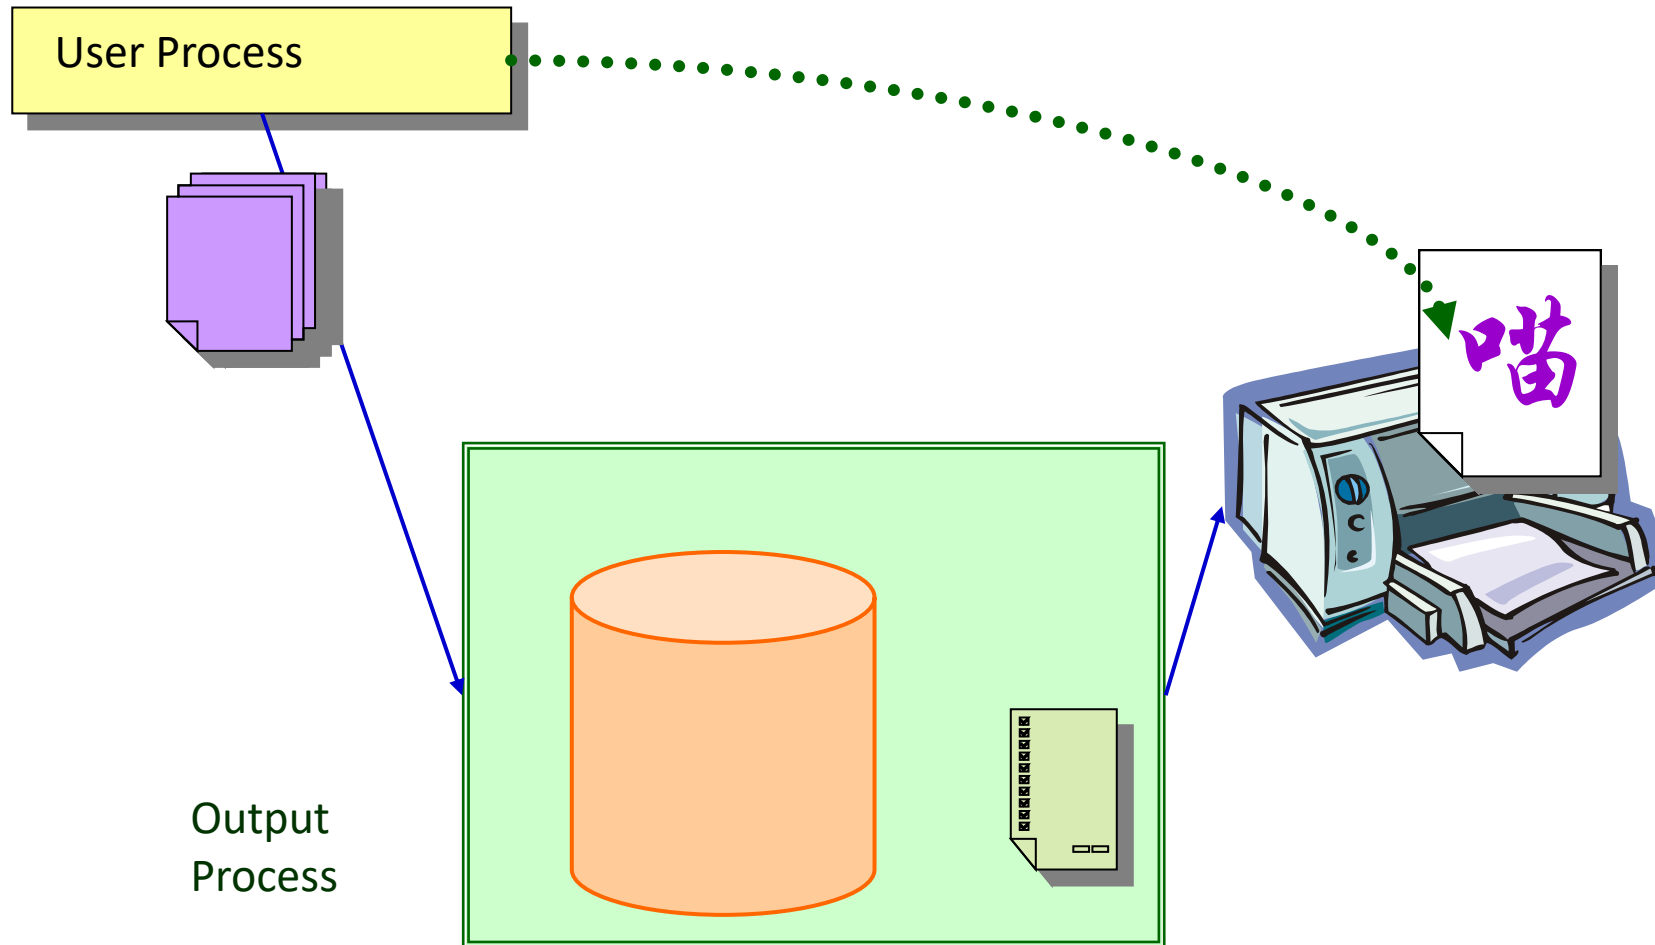
Circular Buffer



Device Management

- **SPOOLing** (Simultaneous Peripheral Operation On Line) hold output for a device
 - If device can serve only one request at a time
 - i.e., Printing

Spooling





任务管理器

文件(F) 选项(O) 查看(V)

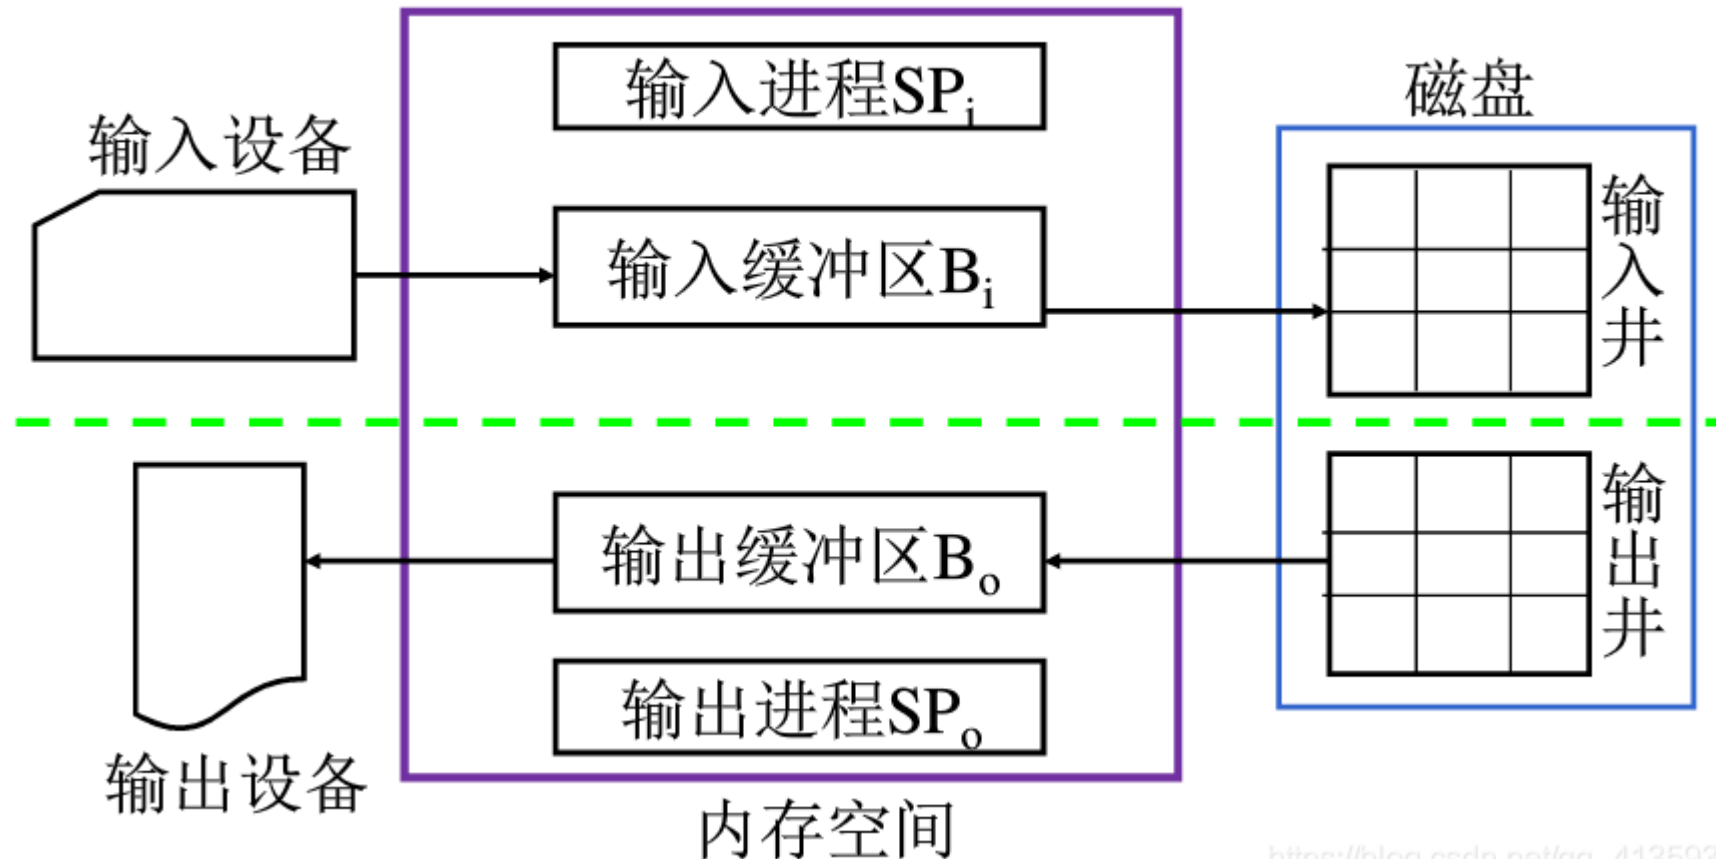
进程 性能 应用历史记录 启动 用户 详细信息 服务

名称	PID	状态	用户名	CPU	内存(活动...	UAC 虚拟化
svchost.exe	6396	正在运行	SYSTEM	00	424 K	不允许
svchost.exe	7172	正在运行	SYSTEM	00	652 K	不允许
svchost.exe	6360	正在运行	dizhang	00	308 K	已禁用
svchost.exe	9340	正在运行	SYSTEM	00	2,268 K	不允许
svchost.exe	13048	正在运行	LOCAL SE...	00	976 K	不允许
svchost.exe	14584	正在运行	LOCAL SE...	00	428 K	不允许
svchost.exe	240	正在运行	SYSTEM	00	628 K	不允许
svchost.exe	16444	正在运行	SYSTEM	00	836 K	不允许
svchost.exe	1380	正在运行	LOCAL SE...	00	184 K	不允许
svchost.exe	4776	正在运行	SYSTEM	00	2,452 K	不允许
svchost.exe	16036	正在运行	SYSTEM	00	1,184 K	不允许
StartMenuExpen...	8348	正在运行	dizhang	00	21,236 K	已禁用
spoolsv.exe	508	正在运行	SYSTEM	00	1,120 K	不允许
SogouImeBroker...	11204	正在运行	dizhang	00	512 K	已禁用
SogouCloud.exe	1708	正在运行	dizhang	00	2,900 K	已禁用
SnippingTool.exe	17124	正在运行	dizhang	00	8,028 K	已禁用
smss.exe	388	正在运行	SYSTEM	00	136 K	不允许
smartscreen.exe	8916	正在运行	dizhang	00	5,616 K	已禁用
SkypeBackground...	8828	已挂起	dizhang	00	0 K	已禁用
SkypeApp.exe	6092	已挂起	dizhang	00	0 K	已禁用
...

简略信息(D) 结束任务(E)

Spooling

设备分为独占式设备，共享使用设备和虚拟设备（即为以spooling使用的外部设备）。独占设备申请后只有到被释放才能被其他进程申请使用，为了让独占设备能**逻辑上**像共享设备一样使用，是一种将独占式设备改造成共享设备的技术（逻辑上）。



Spooling

- SPOOLING技术 (Simultaneous Peripheral Operating On Line)同时联机外围操作技术，它是关于慢速字符设备如何与计算机主机进行数据交换的一种技术，通常又称假脱机技术。在多道程序环境下，利用多道程序中的一道或者两道程序来模拟脱机输入/输出中的外围控制机的功能，以达到“脱机”输入/输出的目的。利用这种技术可把独占设备转变成共享的虚拟设备，从而提高独占设备的利用率和进程的推进速度。
- SPOOLING系统 是对脱机输入/输出工作的模拟，它必须有大容量的且可随机存取的存储器的支持。其主要思想是在联机的条件下，进行两个方向的操作，在数据输入时，将数据从输入设备传送到磁盘或磁带（块设备），然后把这些块设备与主机相连；反过来，在数据输出时，将输出数据传送到磁盘或磁带上，再从磁盘或磁带传送到输出设备。这样，可以将一台独占的物理设备虚拟为并行使用的多态逻辑设备，从而使该物理设备被多个进程共享。

Spooling

- 输入进程SPI是模拟脱机输入时的外围控制机，它将用户要求处理的数据从输入设备通过输入缓冲区再送到输入井（磁盘上开辟的一块区域），当CPU处理这些数据数据时，就直接从输入井读入内存。输出进程SPO是模拟脱机输出时的外围控制机，把用户要求输出的数据，先从内存送到输出井，待输出设备空闲时，再将输出井中的数据通过输出缓冲区（内存中一块区域）传送到输出设备上。

Spooling

- 实例——利用打印机实现打印机共享已经被广泛用于多用户系统和计算机网络中，它实际上就是利用SPOOLING技术将独占的打印机改造为一台供多个用户共享的设备，只要有足够的外存空间和多道程序操作系统的支持即可。
- 1、当用户进程请求打印输出时，SPOOLING系统立即同意为该进程执行打印输出，但并不是真正地把打印机分配给该用户进程，而只是为该进程做两项工作：一项是由输出进程SPO在输出井中为之申请一个空闲的存储空间，并将要打印的数据传送其中存放；另一项工作就是由输出进程SPO再为用户进程申请一张空白的用户请求打印表，并将用户的打印请求填入其中，然后将该表挂到打印机的请求队列上。这时，如果还有另一个进程请求打印机时，则系统仍同意为该进程执行打印输出，当然，系统所做的工作仍是以上两项内容。
 - 2、在打印机执行实际打印时，如果打印机空闲，输出进程SPO将从请求打印队列的队首取出一张打印表，根据打印表中的要求将要打印的数据从输出井传送到内存输出缓冲区，再传送到打印机打印。打印完后，输出进程SPO将再检查请求打印队列中是否还有待打印的请求表，若有则再取出一张请求打印表，将新的但因要求继续打印。如此反复，直到请求打印队列空为止，输出进程才将自己阻塞起来，并在下次再有打印请求时被唤醒。



At last

Layered I/O software system

user-space I/O software

device-independent I/O software

device driver

interrupt handlers

hardware



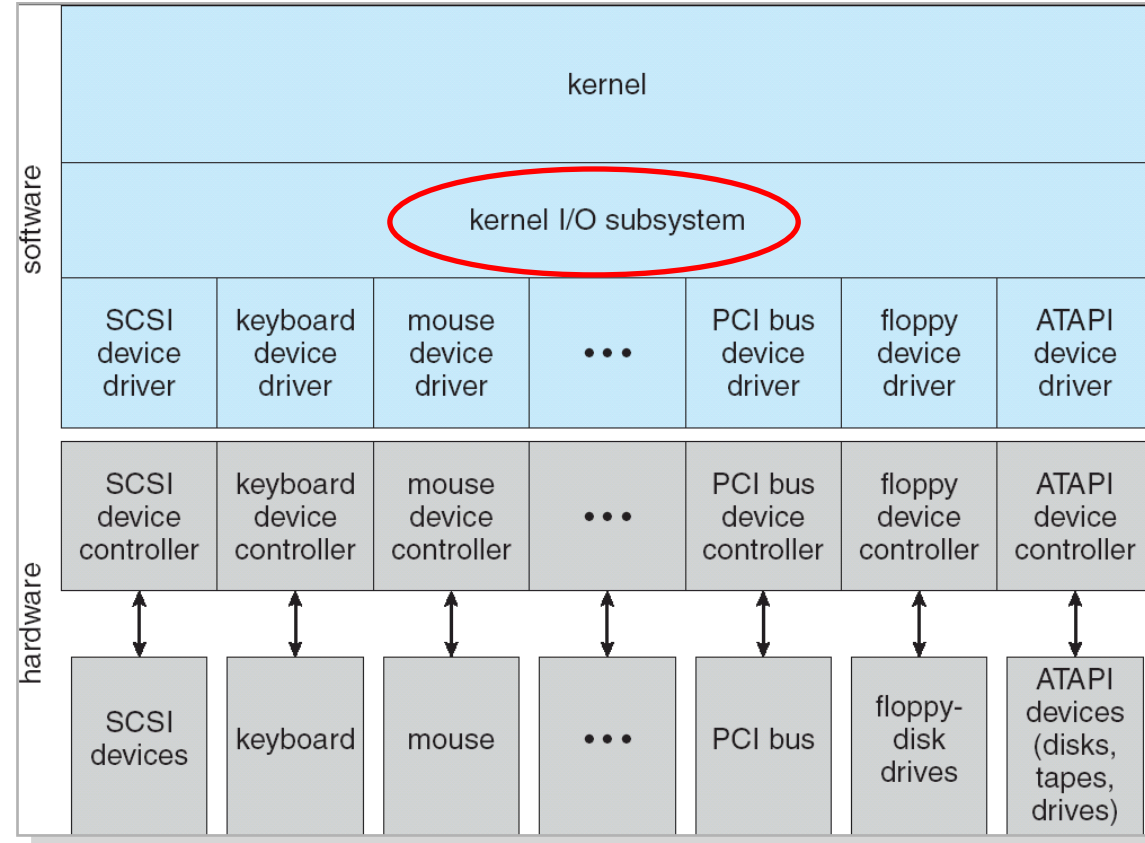
Summary

Summary

- Thousands of devices, varying in many dimensions
 - **Character-stream** or **block**
 - **Sequential** or **random-access**
 - **Sharable** or **dedicated**
 - **Speed of operation**
 - **read-write, read only, or write only**
- In common
 - **Port**
 - **Bus** (**daisy chain** or shared direct access)
 - **Controller**
 - I/O instructions control devices
 - Devices have addresses

Summary

- Hides differences from kernel
 - Done by Kernel I/O Subsystem



Three Main Problems

- Thousands of devices, each slightly different
 - **Provide Uniform Interfaces, Despite Wide Range of Different Devices**
- Devices unreliable: media failures and transmission errors
- Devices unpredictable and/or slow
 - **Devices Controlling and Management**

Summary

- Device management
 - Scheduling
 - Buffering
 - Error handling
- Spooling



北京交通大学

Thank you!

Q & A

