

Operating system

Part VII: Deadlock [死锁]

By KONG LingBo (孔令波)

Outline of topics covered by this course

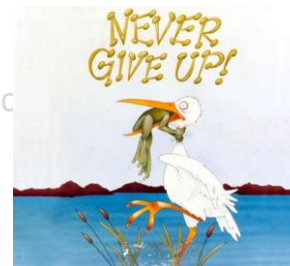
- **INTRODUCTION** – why should we learn OS?
- **OVERVIEW** – what problems are considered by modern OS in more details?
- **EXECUTION** – CPU management
 - Process and Thread
 - CPU scheduling
- **EXECUTION** – competition (synchronization problem)
 - Synchronization
 - Deadlock

- **Deadlock**
 - **Definition, Model**
 - Four necessary conditions
- **Methods for Handling Deadlocks**
 - Providing **enough resources**
 - Staying Safe
 - **Preventing** Deadlocks
 - **Avoiding** Deadlocks
 - Living Dangerously
 - Let the deadlock happen, then **detect** it and **recover** from it.
 - **Ignore** the risks

However, it also will cause some problems if

- If there is **no controlled access** to shared data, execution of the processes on these data can interleave ← **Cooperation**.
 - The results will then depend on the order in which data is modified → **Data Inconsistency**
 - i.e. the results are non-deterministic
- Concurrent processes (or threads) often need to **share data** (maintained either in shared memory or files) and **resources**
 - If there is no proper policy to assign resources among processes, it may result in that all the processes get blocked → **Deadlock** [死锁]

We have discussed this by **synchronization**





Deadlock: General

From Wikipedia, the free encyclopedia

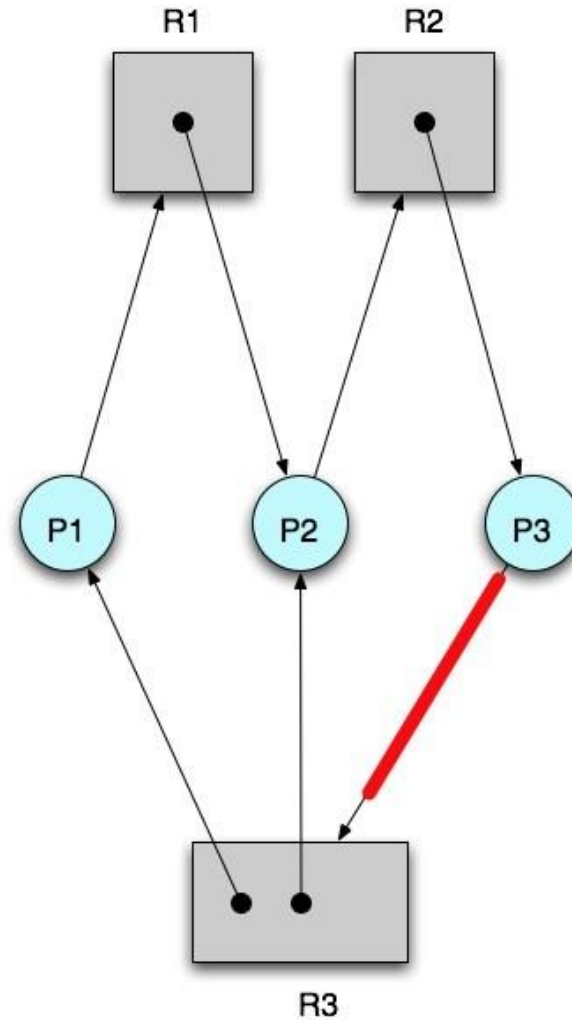
- A deadlock is a situation wherein **two or more competing actions are waiting for the other to finish, and thus neither ever does.**
 - It is often seen in a paradox like the "chicken or the egg."



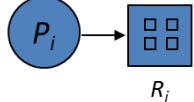
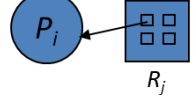
“ When two trains approach each other at a crossing, both shall come to a full stop and neither shall start up again until the other has gone. ”

— Illogical [statute](#) passed by the [Kansas Legislature](#)

Deadlock in OS

- A set of blocked processes each holding some resources and waiting to acquire a resource held by another process in the set.
- None of the processes can proceed or back-off (release resources it owns)



- Process 
- Resource Type with 4 instances 
- P_i requests instance of R_j 
- P_i is holding an instance of R_j 

It seems ... [Basic Facts, Theorem]

1. If there are enough materials for everyone, no deadlock at all!
2. When a deadlock occurs, four conditions must have been reached at the same time!

① Mutual Exclusion [互斥]

- At least one resource is non-sharable: at most one process at a time can use it

② Hold-and-Wait [占有并等待]

- At least one process is holding one resource while waiting to acquire others, that are being held by other processes

Deadlock can arise only if ...

- A deadlock can arise **only if** all four conditions hold

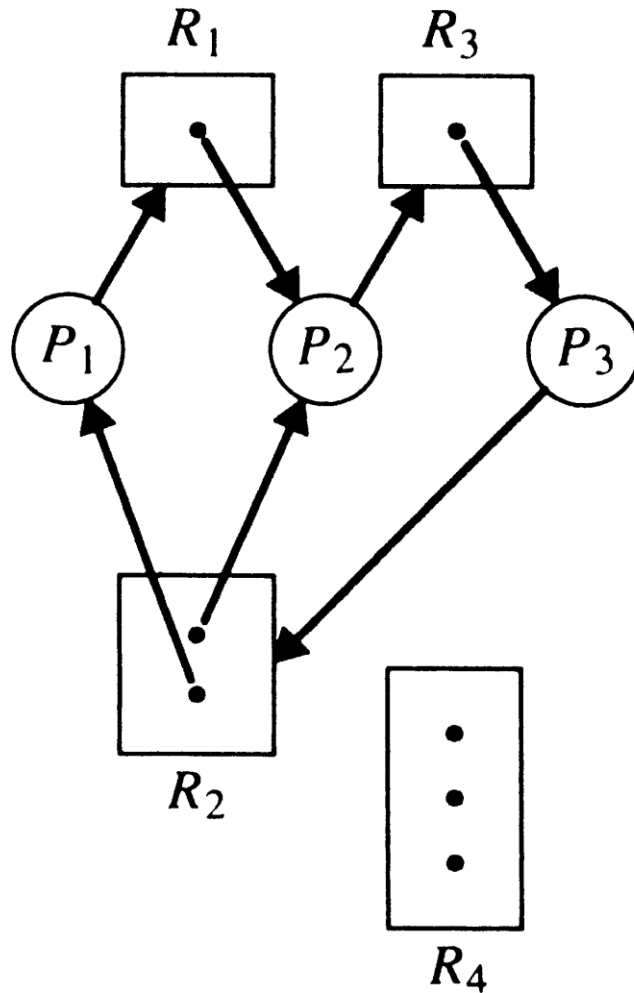
③ No preemption [非抢占]

- A resource cannot be preempted (a process needs to give it up voluntarily)

④ Circular Wait [循环等待]

- There exists a set $\{P_0, P_1, \dots, P_n\}$ of waiting processes such that
 - P_i is waiting for a resource that is held by P_{i+1} , $0 \leq i < n$
 - P_n is waiting for a resource that is held by P_0

Deadlock

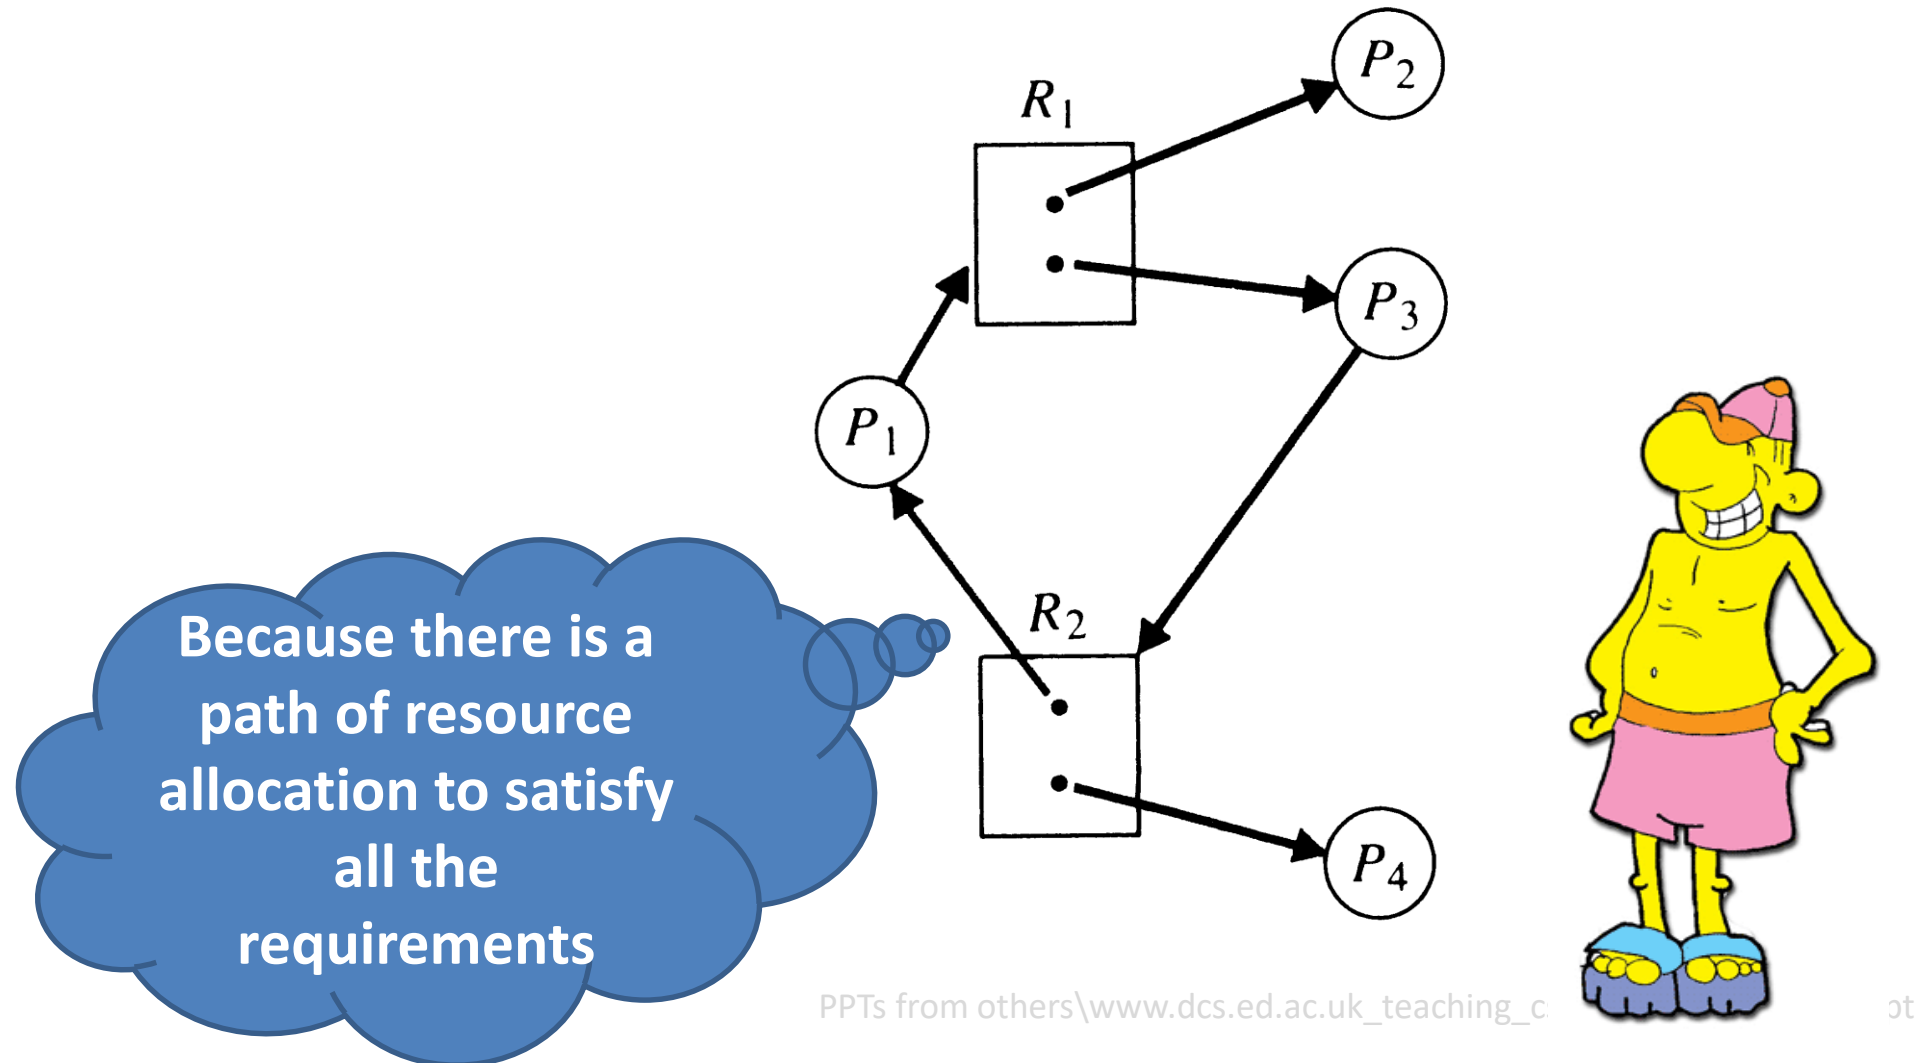


Because you cannot find **an execution sequence** of the all three processes so that each of them can finishes its work finally.



PPTs from others\www.dcs.ed.ac.uk_teaching_cs3_osslidesdeadlock.ppt

A cycle not sufficient to imply a deadlock:



It seems ...

1. If the graph contains no cycles, then no process is deadlocked.
2. If there is a cycle, then two situations:
 - If resource types have multiple instances, then deadlock **MAY** exist.
 - If each resource type has 1 instance, then deadlock has occurred.
 - The existence of a cycle is a sufficient and necessary condition for the existence of a deadlock
 - Each process involved in the cycle is deadlocked

- Deadlock
 - Definition, Model
- Methods for Handling Deadlocks
 - Providing enough resources
 - Staying Safe
 - Preventing Deadlocks
 - Avoiding Deadlocks
 - Living Dangerously
 - Let the deadlock happen, then detect it and recover from it.
 - Ignore the risks

Providing enough resources

- A useful equation!

- Given:
 - Here are **3** processes: A, B, C. Each of them requires **5** system resources.
- Question:
 - How many resources should the system at least have so that the system is safe?
- Rule:
 - If the number of system resources satisfies the following equation, then the system is safe!

$$\sum (P_{\max} - 1) + 1 \leq R_{Total}$$

A useful equation!

$$\sum (P_{\max} - 1) + 1 \leq R_{Total}$$

- P_{\max} : is the **max** number of the required resources by process P
- R_{total} : is the total resources the system has
- It could be simplified as follows, where N is the number of processes

$$(P_{\max} - 1) * N + 1 \leq R_{Total}$$

- It is easy to answer the given question:

$$(P_{\max} - 1) * N + 1 \leq R_{Total}$$

- According to the above equation
 - N is 3
 - P_{\max} is 5
 - So, there should be at least $(5-1)*3+1 = 13$ resources in the system, then the system is in safe!

Variations

- Question:
 - A system has 10 tape drivers, which are shared by m processes, and each process requires 3 tape drivers at most. So, what should “ m ” be then the system could be in safe?
 - A.3 B.4 C.5 D.6
- By that equation, we have
 - $(3-1)*m+1 \leq 10$
 - So, $m \leq 4.5$
 - B is the answer

Staying Safe - Deadlock Prevention (预防)

- Do not allow one of the four conditions to occur.
 - **Mutual Exclusion** [互斥]
 - Only one process may use a resource at a time
 - **Hold and Wait** [持有和等待]
 - A process may hold allocated resources while awaiting assignment of others
 - **No Preemption** [非抢占]
 - No resource can be forcibly removed from a process holding it
 - **Circular Wait** [循环等待]
 - A closed chain of processes exists, such that each process holds at least one resource needed by the next process in the chain

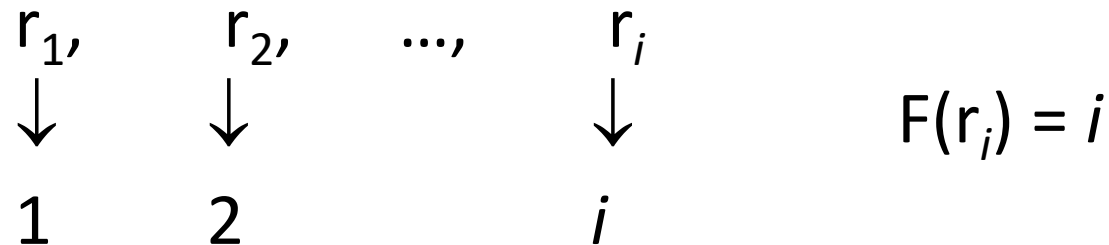
Deadlock Prevention - **negating Hold and Wait**

- Two strategies
 1. When the process begins, all the resources required by it should be assigned to it
 - inefficient - not all resources needed all the time
 - processes probably will not know in advance what resources they will need
 - may have to wait excessive time to get all resources at once - starvation
 - high priority processes may cause starvation of low priority processes

- Two strategies
 2. When a new request is needed by a processes
 - A.it could release existing resources it holds if it fails to get a new resource immediately (try again later)
 - B.the process always releases its existing resources and asks for all of them at once

Deadlock Prevention - **negating Circular Wait**

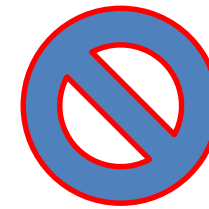
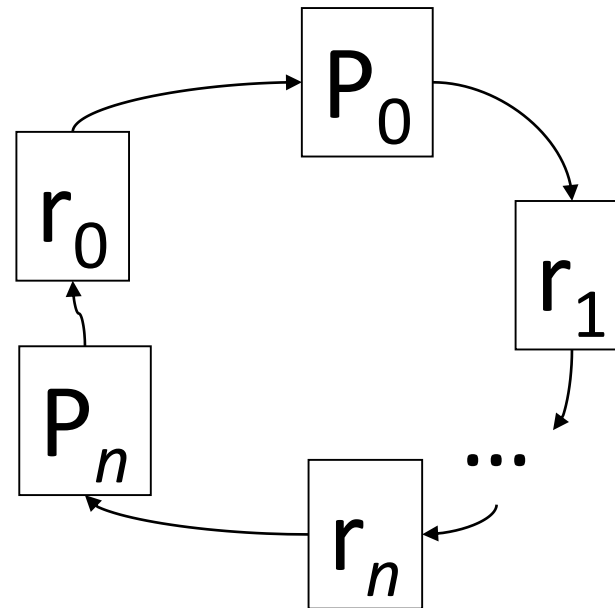
- Ordered resource allocation [资源顺序分配法]
 - Quite popular in management science (such as MBA)
 - Steps:
 - Assign each resource class with a unique number



- The process should apply for all its required m-class resources following the order of resource class numbers

- Since all the resources lower number than the process requires the resource cycle! cycle!

This implies ... request has been known. You've to propose your request at the beginning



NO WAY!

Staying Safe - Deadlock Avoidance (避免)

- Deadlock prevention → low device utilization and reduced system throughput.
- Deadlock avoidance
 - Given the complete sequence of requests and releases for each process, we can decide for each request whether or not the process should wait.
 - For every request, the system
 - considers the resources currently available, the resources currently allocated, and the future (Needed) requests and releases of each process, and
 - decides whether the current request can be satisfied or must wait to avoid a possible future deadlock.

Example 2:

- 5 processes P_0 through P_4 ;
3 resource types:

A (**10** instances), B (**5** instances), and C (**7** instances)

Snapshot at time T_0 :

	<u>Allocation</u>			<u>Max</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
P_0	0	1	0	7	5	3	3	3	2
P_1	2	0	0	3	2	2			
P_2	3	0	2	9	0	2			
P_3	2	1	1	2	2	2			
P_4	0	0	2	4	3	3			

Example 2 (cont')

- We can compute the matrix **Need** as *Max – Allocation*

	<u>Need</u>			
	A	B	C	
P_0	7	4	3	
P_1	1	2	2	
P_2	6	0	0	
P_3	0	1	1	
P_4	4	3	1	

	<u>Available</u>		
	A	B	C
	3	3	2

SAFE or not is
determined by if we could
find a sequence!

Example 2 (cont')

	<u>Need</u>		
	A	B	C
P0	7	4	3
P1	1	2	2
P2	6	0	0
P3	0	0	0
P4	4	3	1

<u>Available</u>		
A	B	C
5	4	3

- Compare the available resources and the needed resources for each process to find if there is some process whose requirement could be satisfied.

– P1 1 2 2

– P3 0 1 1

- Randomly select one, here we select P3

- The available now is

	<u>Allocation</u>		
P3	2	1	1

– $\langle 3, 3, 2 \rangle + \langle 2, 1, 1 \rangle \rightarrow \langle 5, 4, 3 \rangle$

P3 → **P1**

Example 2 (cont')

	<u>Need</u>
	A B C
P0	7 4 3
P1	0 0 0
P2	6 0 0
P3	0 0 0
P4	4 3 1
	<u>Available</u>
	A B C
	3 4 3

- Similarly, there are two processes whose requirements could be satisfied.
 - P1 1 2 2
 - P4 4 3 1
- Randomly select one, here we select P1
- The available now is

	<u>Allocation</u>
P1	2 0 0

 - $\langle 5, 4, 3 \rangle + \langle 2, 0, 0 \rangle \rightarrow \langle 7, 4, 3 \rangle$

P3 → P1 → P0 → P2 → P4

Example 2 (cont')

	<u>Need</u>
	A B C
P0	7 4 3
P1	0 0 0
P2	6 0 0
P3	0 0 0
P4	4 3 1
	<u>Available</u>
	A B C
	7 4 3

- Similarly, there are three processes whose requirements could be satisfied.

– P0 7 4 3

– P2 6 0 0

– P4 4 3 1

- We can directly execute processes in the order:

– P0 → P2 → P4

So, system is
safe at the
snapshot time
t0!

Example 2:

- How about the snapshot as follows?

	<u>Allocation</u>		<u>Need</u>	<u>Available</u>
	A B C		A B C	A B C
P0	0 1 0	P0	7 4 3	1 1 1
P1	2 0 0	P1	1 2 2	
P2	3 0 2	P2	6 0 0	
P3	2 1 1	P3	0 1 1	
P4	0 0 2	P4	4 3 1	

P3

Example 2 (cont')

	<u>Need</u>
	A B C
P0	7 4 3
P1	1 2 2
P2	6 0 0
P3	0 0 0
P4	4 3 1

<u>Available</u>
A B C
3 2 2

- Compare the available resources and the needed resources for each process to find if there is some process whose requirement could be satisfied.

– P3 0 1 1

- So we select P3
- The available now is

– $\langle 1, 1, 1 \rangle + \langle 2, 1, 1 \rangle \rightarrow \langle 3, 2, 2 \rangle$

	<u>Allocation</u>
P3	2 1 1

P3 → P1

Example 2 (cont')

	<u>Need</u>
	A B C
P0	7 4 3
P1	1 2 2
P2	6 0 0
P3	0 0 0
P4	4 3 1

Available
A B C
5 2 2

- Similarly, we can see P1 could be satisfied.
 - P1 1 2 2
- So we select P1
- The available now is
 - $\langle 3, 2, 2 \rangle + \langle 2, 0, 0 \rangle \rightarrow \langle 5, 2, 2 \rangle$

	<u>Allocation</u>
P1	2 0 0

Example 2 (cont')

	<u>Need</u>
	A B C
P0	7 4 3
P1	0 0 0
P2	6 0 0
P3	0 0 0
P4	4 3 1
	<u>Available</u>
	A B C
	5 2 2

- After P1, available res is $\langle 5, 2, 2 \rangle$
- Sadly, this time, the rest three processes could not be satisfied $\langle 5, 2, 2 \rangle$.
 - P0 7 4 3
 - P2 6 0 0
 - P4 4 3 1
- This means it is not safe for the snapshot of example 2!

Example 3:

- 5 processes P_0 through P_4 ;
3 resource types:

A (**10** instances), B (**5** instances), and C (**7** instances)

Snapshot at time T_0 :

	<u>Allocation</u>			<u>Max</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
P_0	0	1	0	7	5	3	3	3	2
P_1	2	0	0	3	2	2			
P_2	3	0	2	9	0	2			
P_3	2	1	1	2	2	2			
P_4	0	0	2	4	3	3			

Could the request of
 $P_0 = \langle 2 \ 1 \ 1 \rangle$ be satisfied
or not?

- The logic for that satisfaction is
 - Pretend to satisfy the request first, and check if the change will lead the system into unsafe state or not
 - Still safe, the request could be satisfied
- We need update the resource allocation first – the available and the needed for P0

	<u>Allocation</u>	<u>Needed</u>	<u>Available</u>
	A B C	A B C	A B C
P0	2 2 1	5 3 2	1 2 1
P1	2 0 0	1 2 2	
P2	3 0 2	6 0 0	
P3	2 1 1	0 1 1	
P4	0 0 2	4 3 1	

You've learned how to verify if the system is safe or not from Example 1

Drawbacks of Banker's Algorithm

- processes are rarely known in advance how many resources they will need
- the number of processes changes as time progresses
- resources once available can disappear
- the algorithm assumes processes will return their resources within a reasonable time
- processes may only get their resources after an arbitrarily long delay
- **practical use is therefore rare!**

Living Dangerously -

- the Ostrich[鸵鸟] or **Head-in-the-Sand** algorithm

D.J.['ɔstrɪtʃ]

- **Of course**, Try to reduce chance of deadlock as far as reasonable
- **And**, accept that deadlocks will occur occasionally
 - example: kernel table sizes - max number of pages, open files etc.
- **Because**, maybe
 - **MTBF** versus deadlock probability ?
 - cost of any other strategy may be too high
 - overheads and efficiency

MTBF : mean-time between "failures"

Most Operating systems do this!!



Maybe it's better

— “**Don't hide** from the fact..., Be alert!”



When deadlock happened - Detect & Recover

- **Check** for deadlock (periodically or sporadically[偶发地,零星地]), then **recover**
- Differentiate between
 - Serially reusable resources: A unit must be allocated before being released
 - Consumable resources: Never release acquired resources; resource count is the number currently available

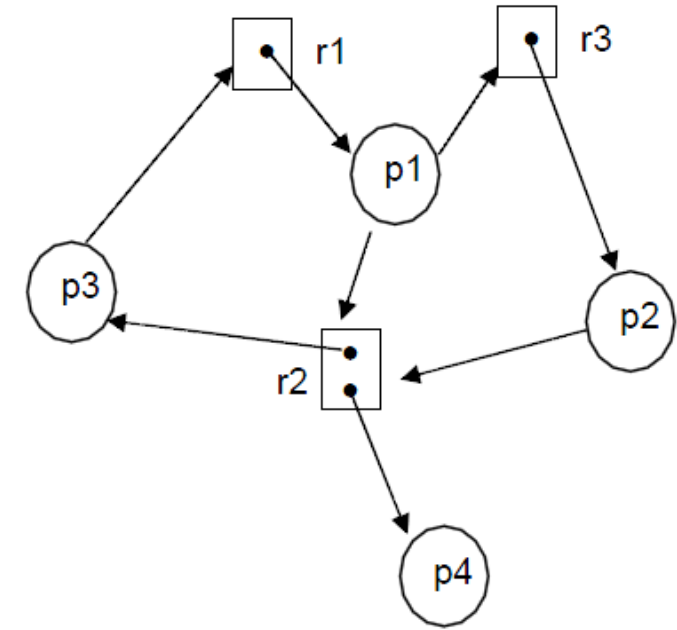
Borrow Banker's algorithm for several instances

Available = [0 0 0]

$$\text{Allocation} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \\ 0 & 0 & 1 \end{bmatrix}$$

$$\text{Request} = \begin{bmatrix} 0 & 1 & 1 \\ 0 & 0 & 1 \\ 1 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

$$\text{Finish} = \begin{bmatrix} \textit{False} \\ \textit{False} \\ \textit{False} \\ \textit{False} \end{bmatrix}$$



- Find if there is a sequence of the involved processes to finish
 - If all the values in Finish vector are TRUE or not

Recovery

- Two strategies

1. Process Termination

- Abort all deadlocked processes.
- Abort one process at a time until the deadlock cycle is eliminated.

– In which order should we choose to abort?

- Priority of the process.
- How long process has computed, and how much longer to completion.
- Resources the process has used.
- Resources process needs to complete.
- How many processes will need to be terminated.
- Is process interactive or batch?

Recovery

- Two strategies

2. Resource Preemption

- Choose a blocked process
- Preempt it (releasing its resources)
 - Back up each deadlocked process to some previously defined checkpoint
- Run the detection algorithm
- Iterate it until the state is not a deadlock state
- Selection Criteria Deadlocked Processes
 - Least amount of processor time consumed so far
 - Least number of lines of output produced so far
 - Most estimated time remaining
 - Least total resources allocated so far
 - Lowest priority

Combined Approach to Deadlock Handling

- Combine the three basic approaches (**prevention**, **avoidance**, and **detection**), allowing the use of the optimal approach for each class of resources in the system.
- Partition resources into hierarchically ordered classes; Use most appropriate technique for handling deadlocks within each class.
- An example:
 - Internal resources (Prevention through resource ordering)
 - Central memory (Prevention through preemption)
 - Job resources (Avoidance)
 - Swappable space (Pre-allocation)

In conclusion

- Deadlock is a situation wherein two or more competing actions are waiting for the other to finish, and thus neither ever does.
- Four necessary conditions
 - Mutual Exclusion [互斥]
 - Hold-and-Wait [占有并等待]
 - No preemption [非抢占]
 - Circular Wait [循环等待]
- Strategies to overcome the deadlock situation
 - Providing enough resources
 - Staying Safe
 - **Preventing** Deadlocks
 - **Avoiding** Deadlocks → **Banker's algorithm!**
 - Living Dangerously
 - Let the deadlock happen, then **detect** it and **recover** from it.

- ICQ B [10 pts]
 - Next week
 - 1 hour, close
 - Blank-filling, Multiple choice, Computation